

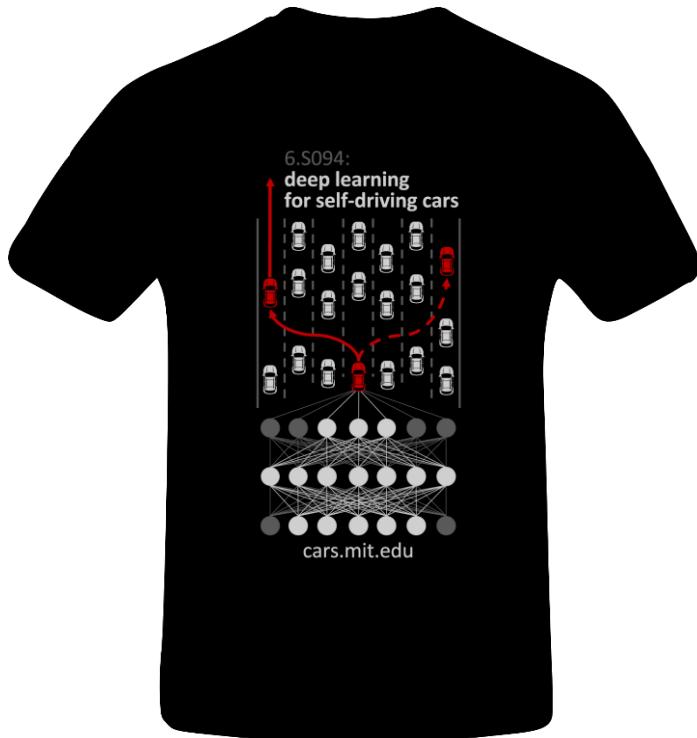


6.S094: Deep Learning for Self-Driving Cars

# Learning to Drive: Convolutional Neural Networks and End-to-End Learning of the Full Driving Tasks

[cars.mit.edu](http://cars.mit.edu)

# Administrative



- **Website:** [cars.mit.edu](http://cars.mit.edu)
- **Contact Email:** [deepcars@mit.edu](mailto:deepcars@mit.edu)
- **Required:**
  - Create an account on the website.
  - Follow the tutorial for each of the 2 projects.
- **Recommended:**
  - Ask questions
  - Win competition!
- **Office hours:** Friday, 5-7pm  
*(more info coming soon)*

# Administrative

- **Website:** [cars.mit.edu](http://cars.mit.edu)
- **Contact Email:** [deepcars@mit.edu](mailto:deepcars@mit.edu)
- **Required:**
  - Create an account on the website.
  - Follow the tutorial for each of the 2 projects.
- **Recommended:**
  - Ask questions
  - Win competition!

# Schedule

Mon, Jan 9	Introduction to Deep Learning and Self Driving Cars
Tue, Jan 10	<b>Learning to Move:</b> Reinforcement Learning for Motion Planning
	<b>DeepTraffic:</b> Solving Traffic with Deep Reinforcement Learning
Wed, Jan 11	<b>Learning to Drive:</b> End-to-End Learning for the Full Driving Task
	<b>DeepTesla:</b> End-to-End Learning from Human and Autopilot Driving
Thu, Jan 12	<b>Karl Iagnemma:</b> From Research to Reality: Testing Self-Driving Cars on Boston Public Roads
Fri, Jan 13	<b>John Leonard:</b> Mapping, Localization, and the Challenge of Autonomous Driving
Tue, Jan 17	<b>Chris Gerdes:</b> TBD
Wed, Jan 18	<b>Sertac Karaman:</b> Past, Present, and Future of Motion Planning in a Complex World
Thu, Jan 19	<b>Learning to Share:</b> Driver State Sensing and Shared Autonomy
Fri, Jan 20	<b>Eric Daimler:</b> The Future of Artificial Intelligence Research and Development
	<b>Learning to Think:</b> The Road Ahead for Human-Centered Artificial Intelligence

# DeepTraffic Leaderboard

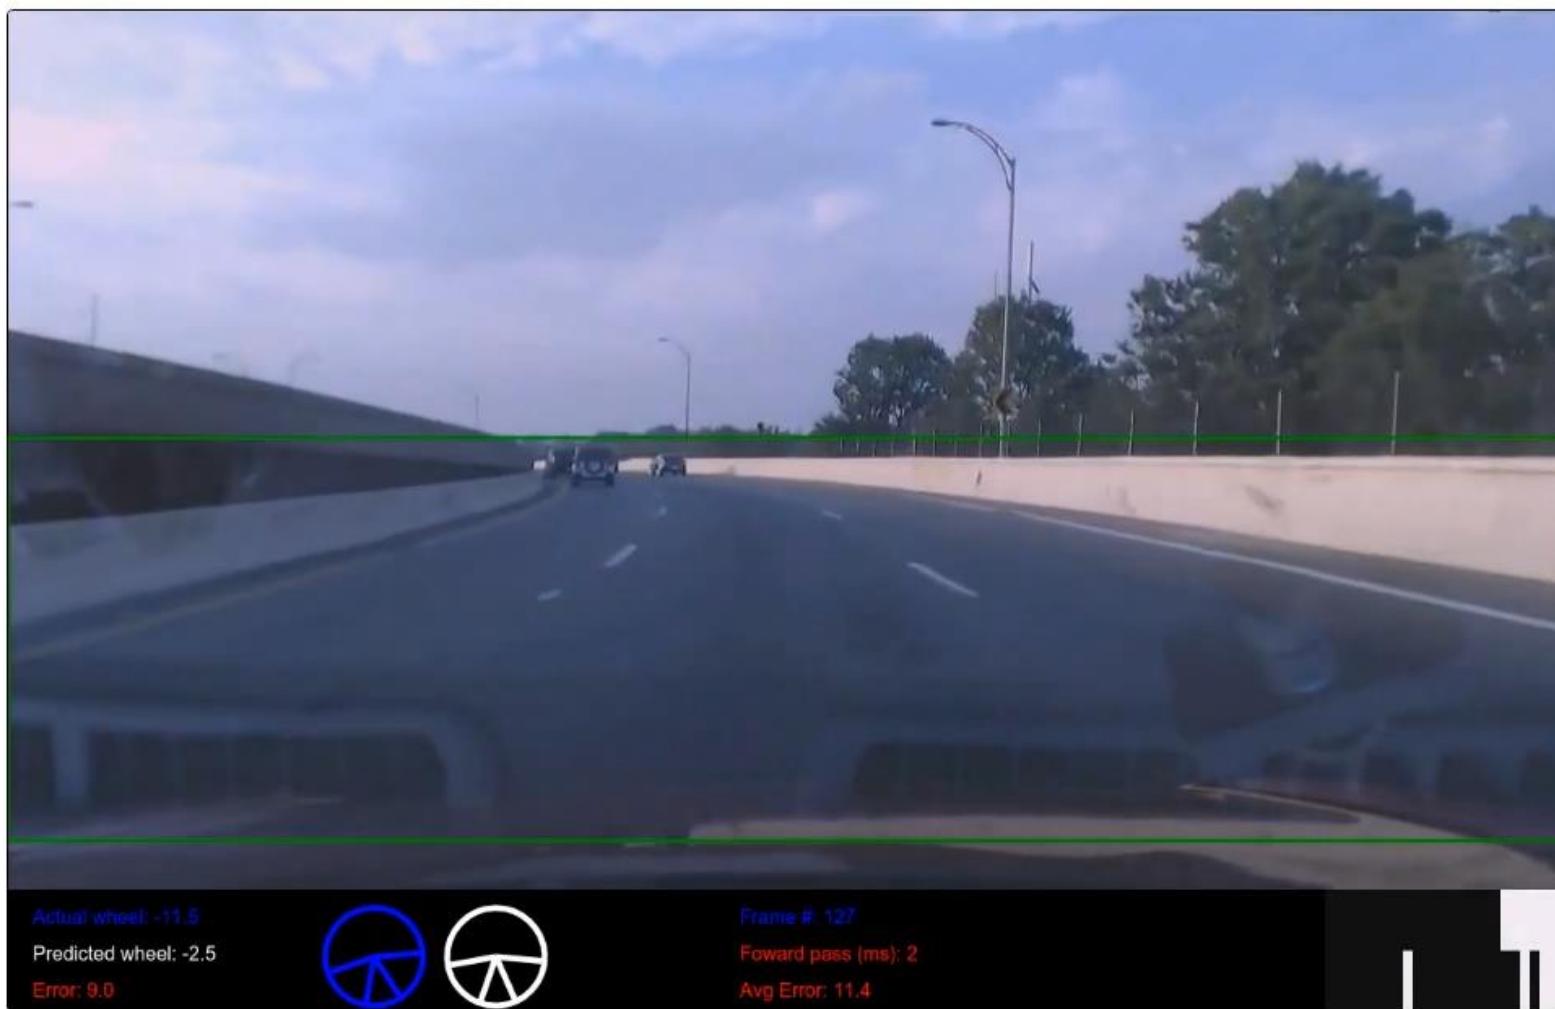


Rank	User	MPH
1	cmauck10	72.67
2	Jeffrey Li	72.45
3	Ted Grunberg	72.31
4	Xiaoxue Wang	72.31
5	Ethan Weber	71.73
6	Indra den Bakker	71.35
7	Kavya R.	71.24
8	Rakesh	71.21
9	LCMartin	71.02
10	anyati	70.95

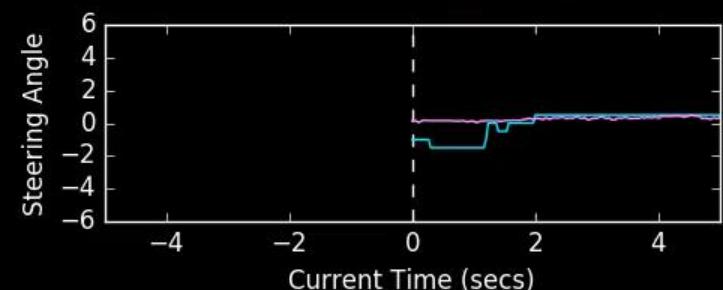
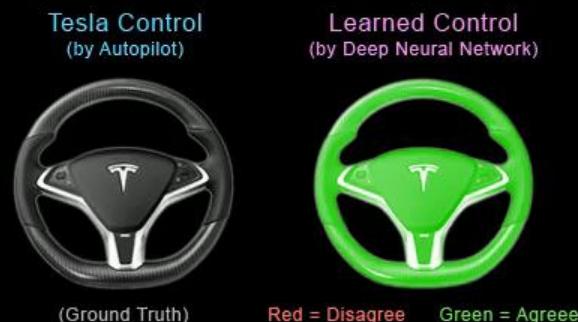
# Illustrative Case Study: Traffic Light Detection



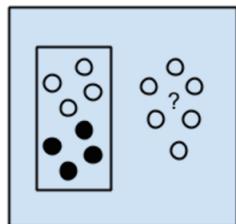
# DeepTesla: End-to-End Learning from Human and Autopilot Driving (in ConvnetJS)



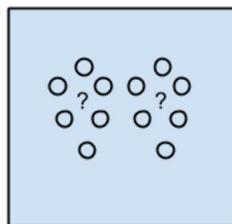
# DeepTesla: End-to-End Learning from Human and Autopilot Driving (in TensorFlow)



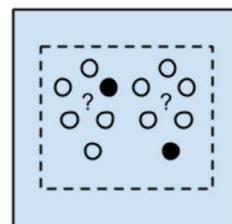
# Computer Vision is Machine Learning



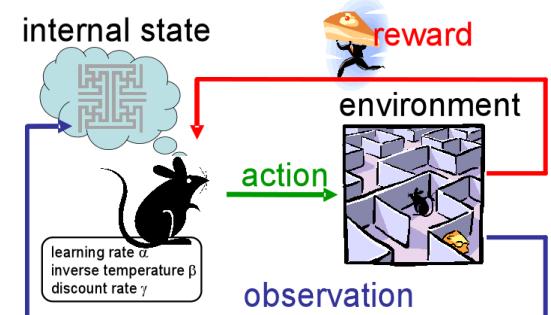
Supervised  
Learning



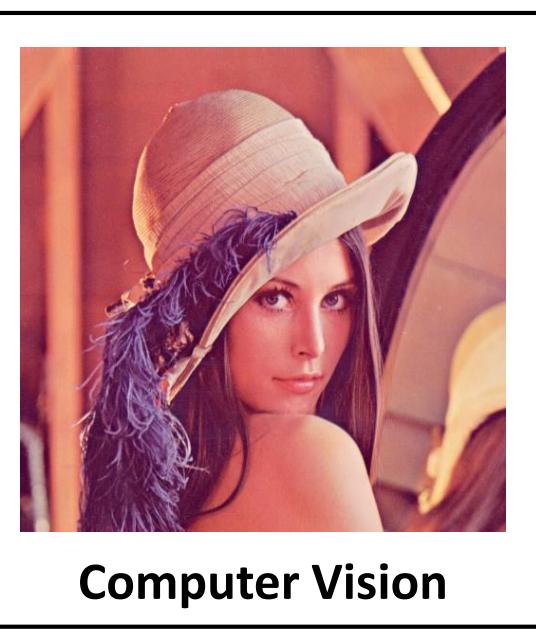
Unsupervised  
Learning



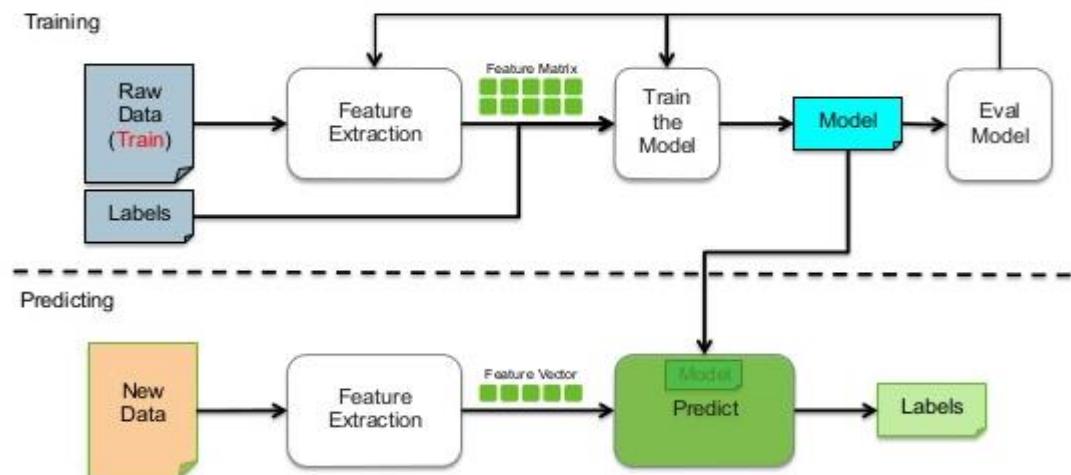
Semi-Supervised  
Learning



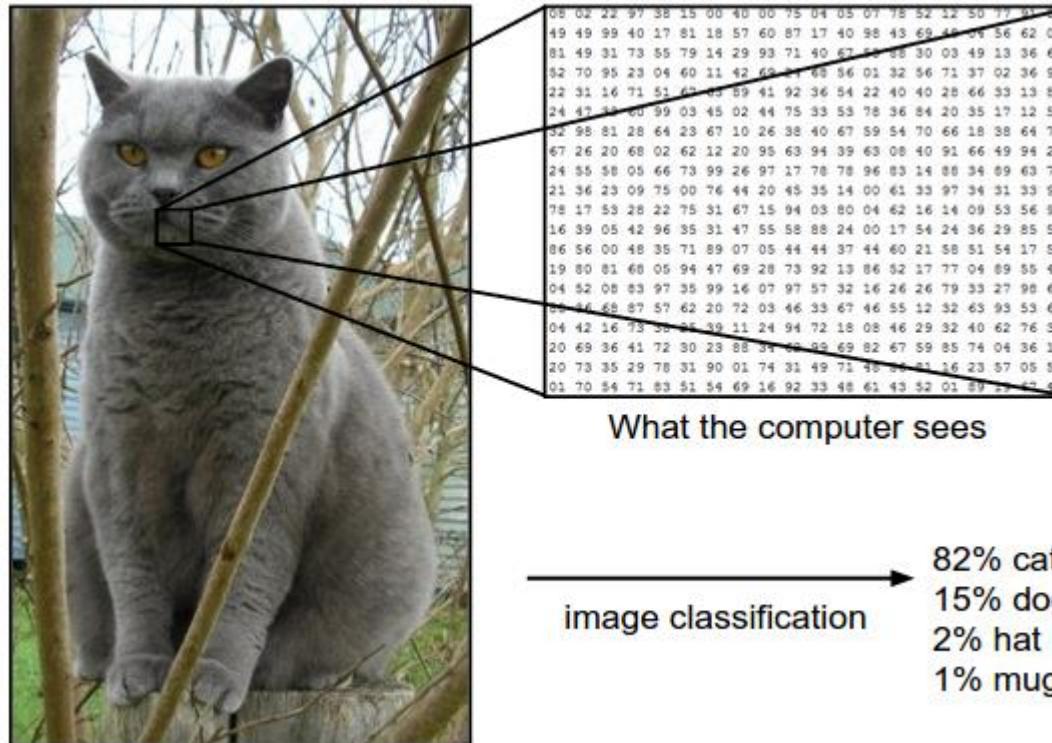
Reinforcement  
Learning



Standard supervised learning pipeline:



# Images are Numbers



- **Regression:** The output variable takes continuous values
- **Classification:** The output variable takes class labels
  - Underneath it may still produce continuous values such as probability of belonging to a particular class.

# Computer Vision is Hard

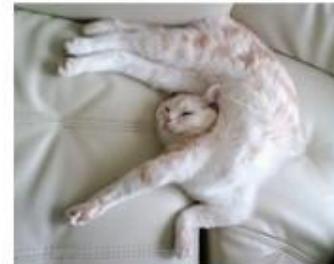
Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



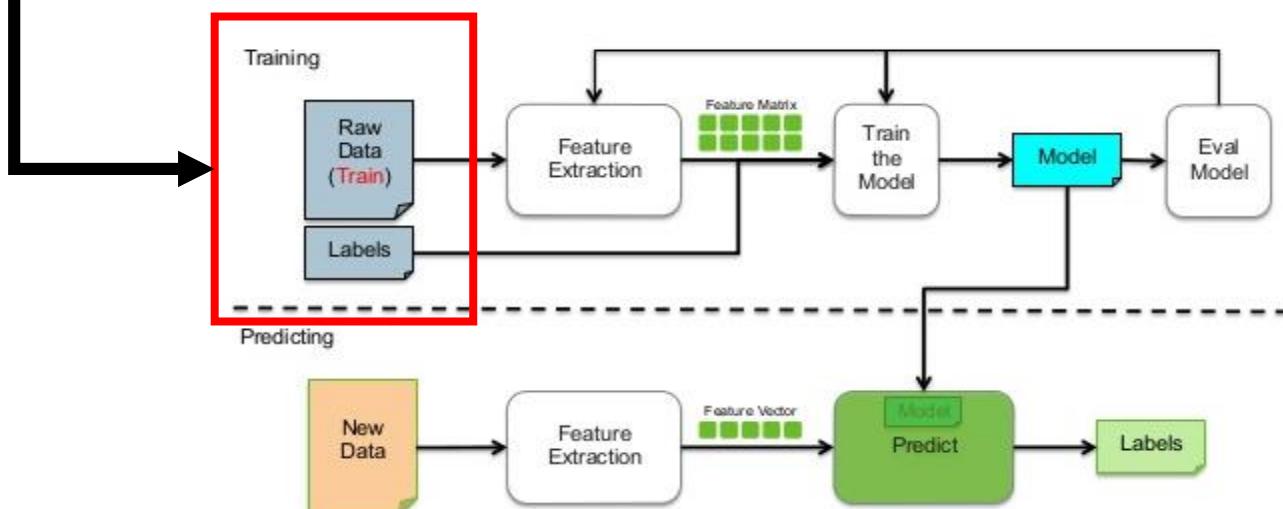
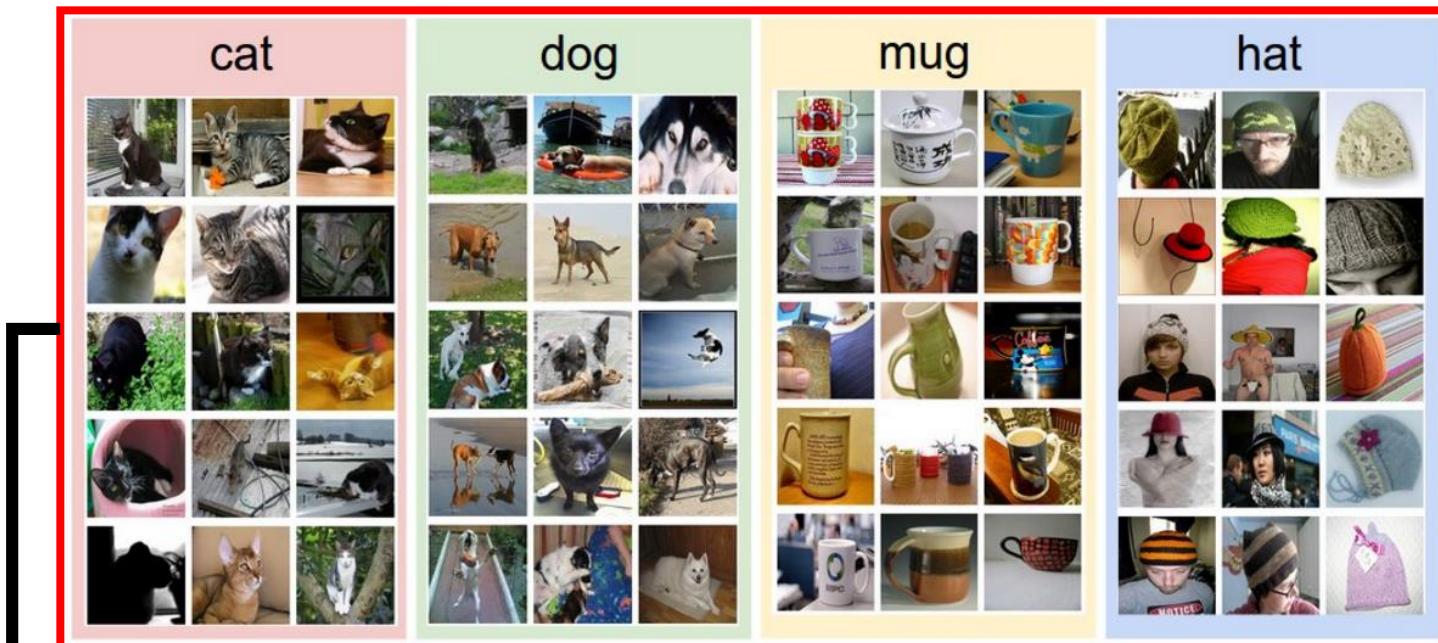
Background clutter



Intra-class variation



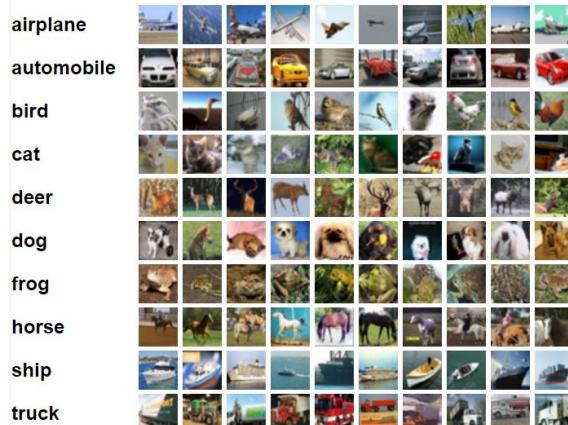
# Image Classification Pipeline



# Famous Computer Vision Datasets

9 3 1 9 7 2 4 5 1 0 3 2 4 3 7 5 9 0 3 4 9  
3 0 2 4 2 9 4 8 3 2 0 1 3 5 3 5 7 4 6 8 5  
2 8 2 3 2 3 8 2 4 9 8 2 9 1 3 9 1 1 1 9 9  
6 3 6 9 0 3 6 0 3 0 1 1 3 9 3 1 5 0 4 9 6  
3 3 8 0 7 0 5 6 9 8 8 4 1 4 4 4 6 9 5 3 3  
1 1 9 5 8 0 4 3 7 7 5 0 5 4 2 0 9 8 1 2 4  
9 5 0 0 5 1 1 1 7 4 7 7 2 6 5 1 8 2 4 1 1  
0 2 1 6 1 7 0 9 5 6 3 2 6 6 7 1 5 2 3 2  
9 4 3 2 1 0 0 2 0 8 7 4 0 9 7 9 3 6 9 3 4  
5 5 1 6 6 2 7 6 7 5 6 6 5 8 1 6 8 7 1 0 5  
7 1 7 5 9 2 3 9 4 3 0 4 5 8 0 0 4 0 4 6 6  
1 6 7 9 6 4 1 1 4 1 3 1 2 3 4 8 1 5 5 0 7  
0 1 6 1 6 7 5 5 5 6 6 8 8 1 7 2 8 3 7 6 5  
6 4 6 8 7 7 1 3 0 7 3 8 6 9 1 6 7 3 6 4 8  
7 9 7 3 1 3 9 7 9 3 6 2 4 9 2 1 4 5 0 3 8

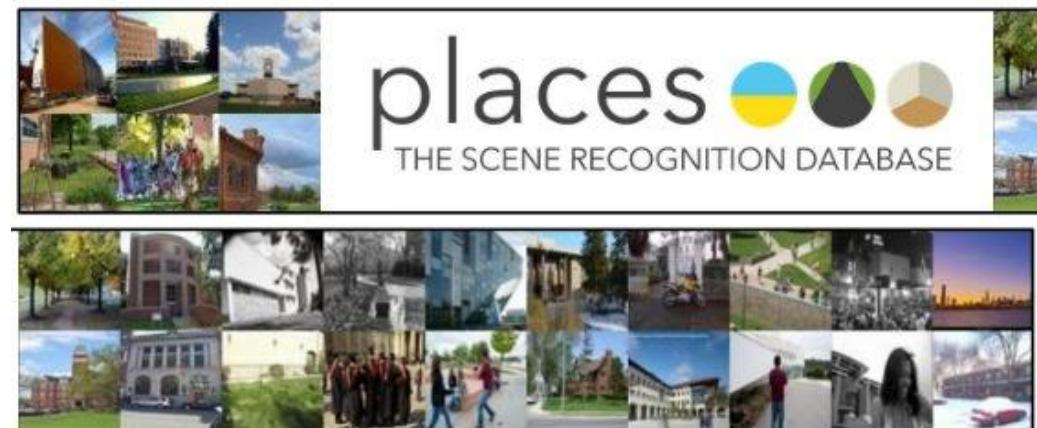
**MNIST:** handwritten digits



**CIFAR-10(0):** tiny images

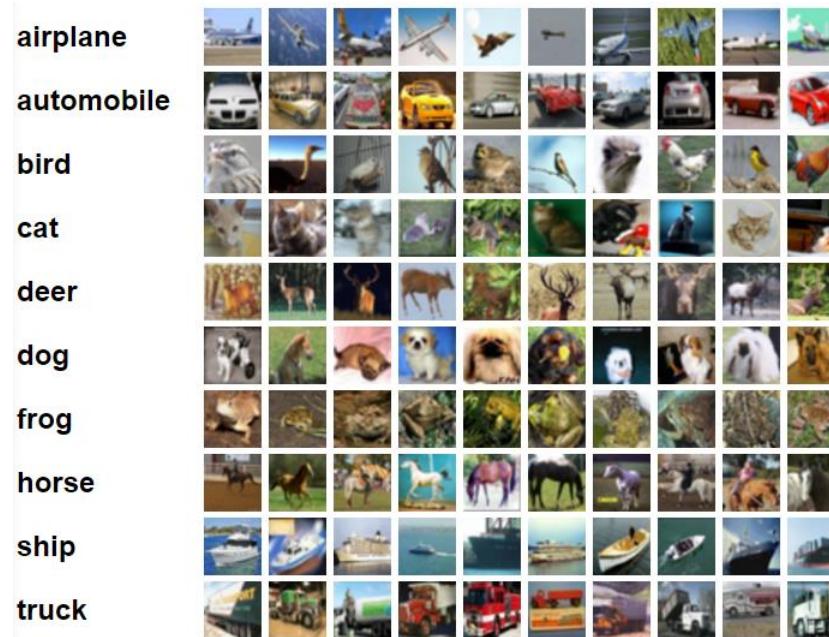
printer housing animal weight drop egg white  
offspring teacher computer headquarters television  
register measure court key structure light date spread  
gallery counter album garage down flower breakfast  
king fireplace church press market lighter  
hotel road paper side site door coffee  
sport screen wall means fan hill can camp fish bathroom  
sky plant house school railcar  
bread wine fox top man car gun  
weapon table man study fly  
cloud cover man net button  
spring range leash van suite mirror seat menu ball flashglass  
descent fruit dog roll sign  
bed shop people sign  
kitchen train kit goal  
engine camera memoriesieve center barwatch  
box stone child overall step  
chain boat tea stand  
apple girl flat student  
home room castle ocean  
flag bank office rule hall  
radio valley cross chair mine t-shirt club  
beach library stage video food building  
baseball material player machine security call clock  
tool football hospital match equipment cell phone mountain telephone  
short circuit bridge scale gas pedal microphone recording crowd

**ImageNet:** WordNet hierarchy



**Places:** natural scenes

# Let's Build an Image Classifier for CIFAR-10



test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

→ 456

# Let's Build an Image Classifier for CIFAR-10

$$\begin{array}{c} \text{test image} \\ \left| \begin{array}{cccc} 56 & 32 & 10 & 18 \\ 90 & 23 & 128 & 133 \\ 24 & 26 & 178 & 200 \\ 2 & 0 & 255 & 220 \end{array} \right| - \begin{array}{c} \text{training image} \\ \left| \begin{array}{cccc} 10 & 20 & 24 & 17 \\ 8 & 10 & 89 & 100 \\ 12 & 16 & 178 & 170 \\ 4 & 32 & 233 & 112 \end{array} \right| = \begin{array}{c} \text{pixel-wise absolute value differences} \\ \left| \begin{array}{cccc} 46 & 12 & 14 & 1 \\ 82 & 13 & 39 & 33 \\ 12 & 10 & 0 & 30 \\ 2 & 32 & 22 & 108 \end{array} \right| \rightarrow 456 \end{array} \end{array} \end{array}$$



## Accuracy

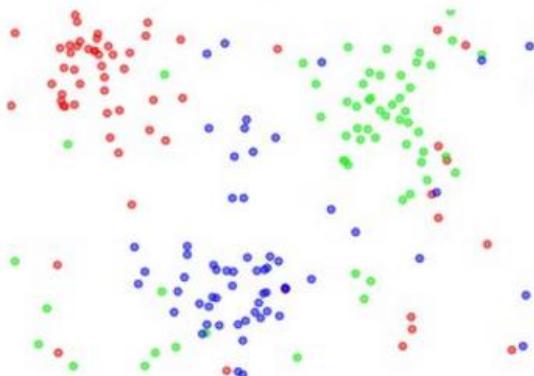
Random: **10%**

Our image-diff (with L1): **38.6%**

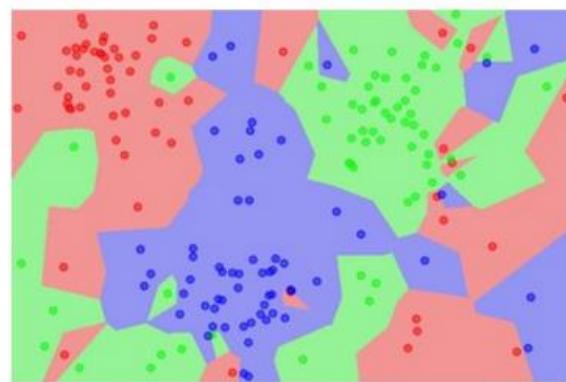
Our image-diff (with L2): **35.4%**

# K-Nearest Neighbors: Generalizing the Image-Diff Classifier

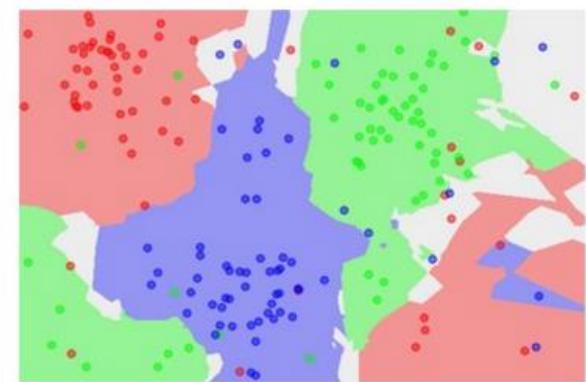
the data



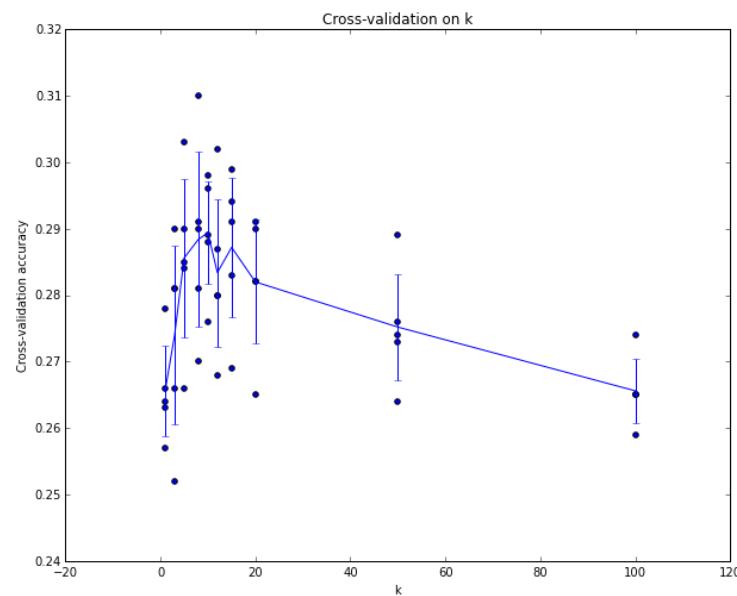
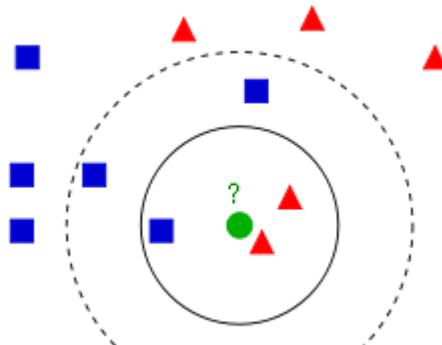
NN classifier



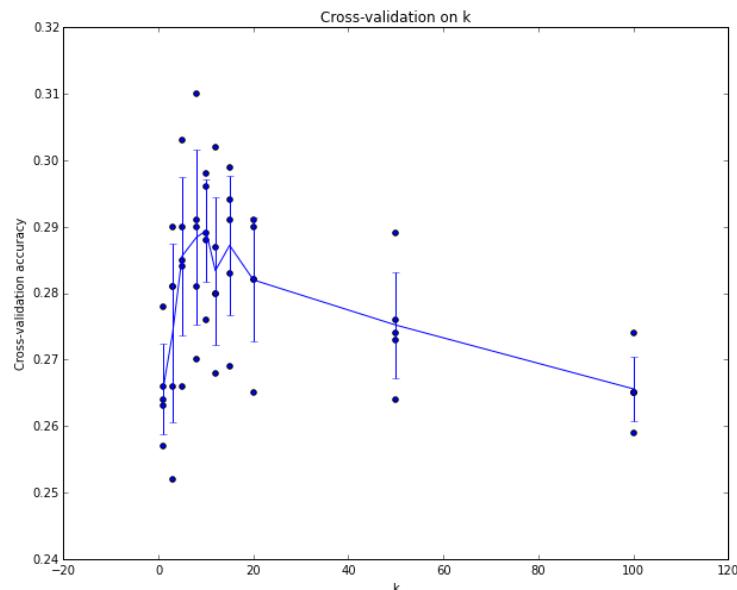
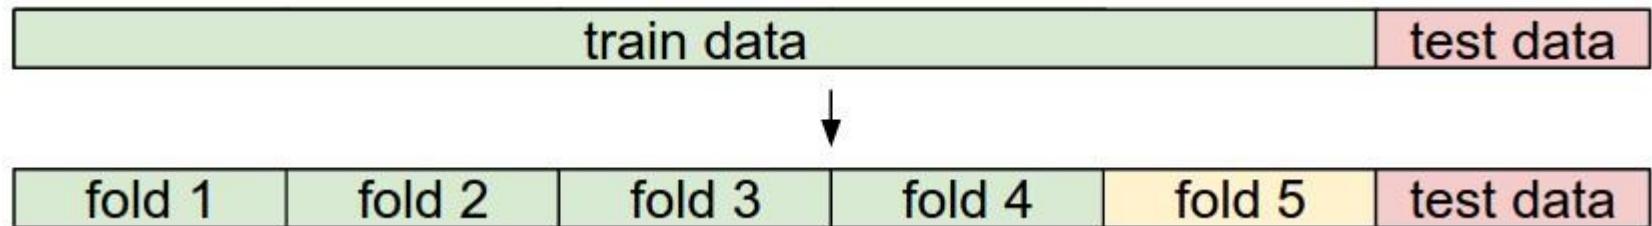
5-NN classifier



Tuning (hyper)parameters:



# K-Nearest Neighbors: Generalizing the Image-Diff Classifier



## Accuracy

Random: **10%**

Training and testing on the same data: **35.4%**

7-Nearest Neighbors: **~30%**

Human: **~94%**

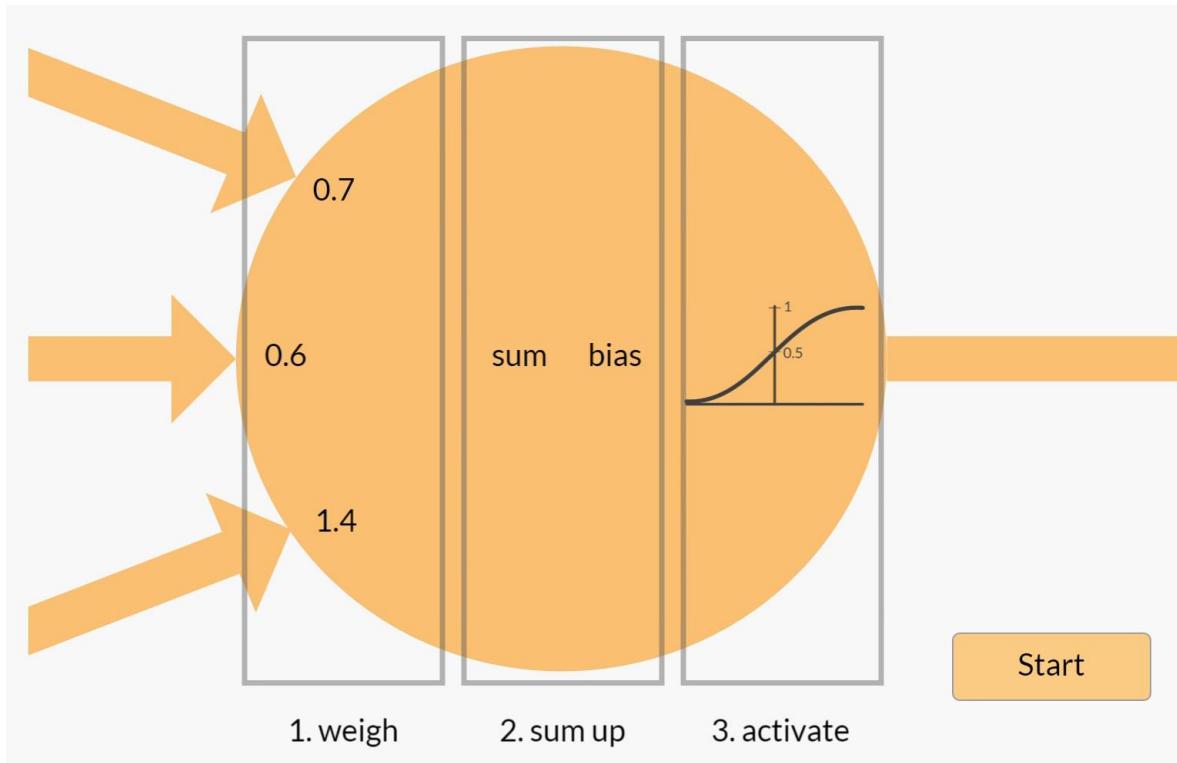
...

Convolutional Neural Networks: **~95%**

# Reminder: Weighing the Evidence

Evidence

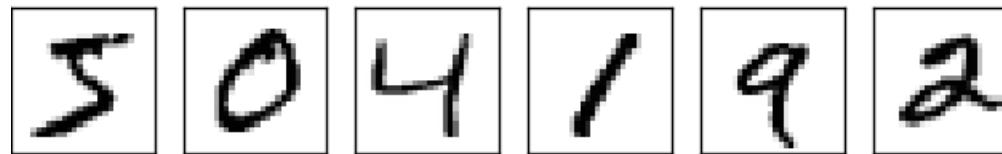
Decisions



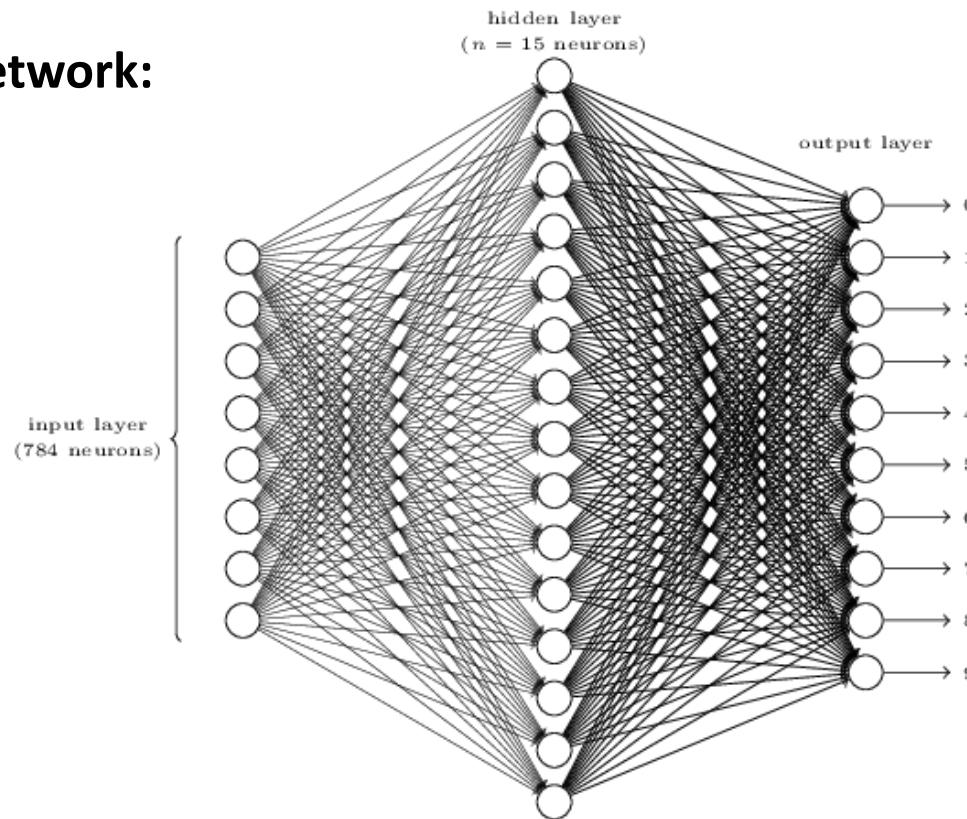
$$\text{output} = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

# Reminder: Classify and Image of a Number

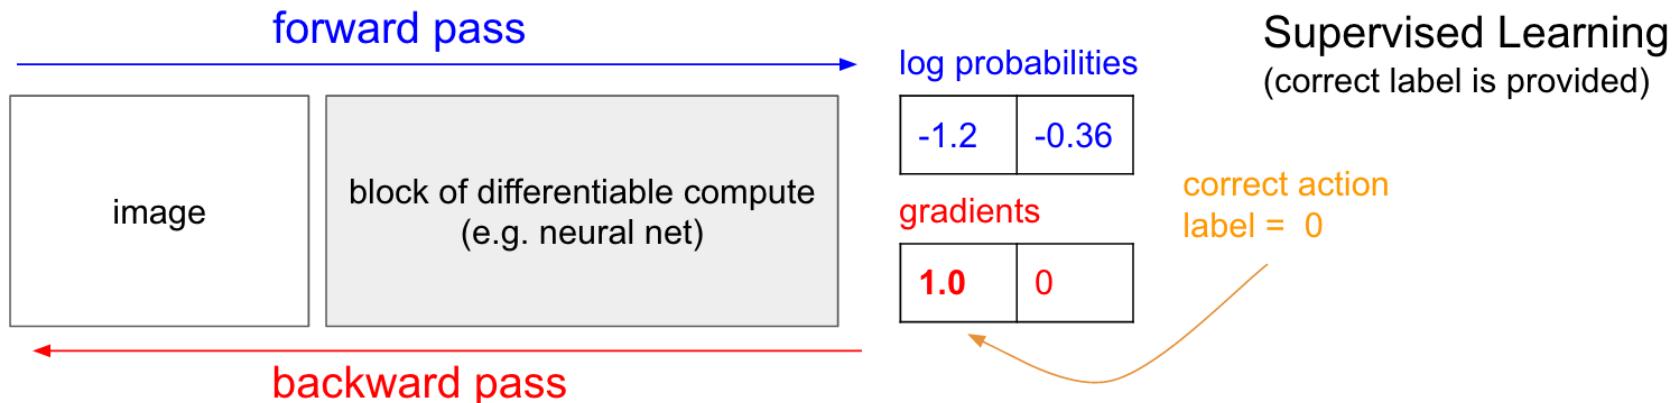
**Input:**  
(28x28)



**Network:**



# Reminder: “Learning” is Optimization of a Function

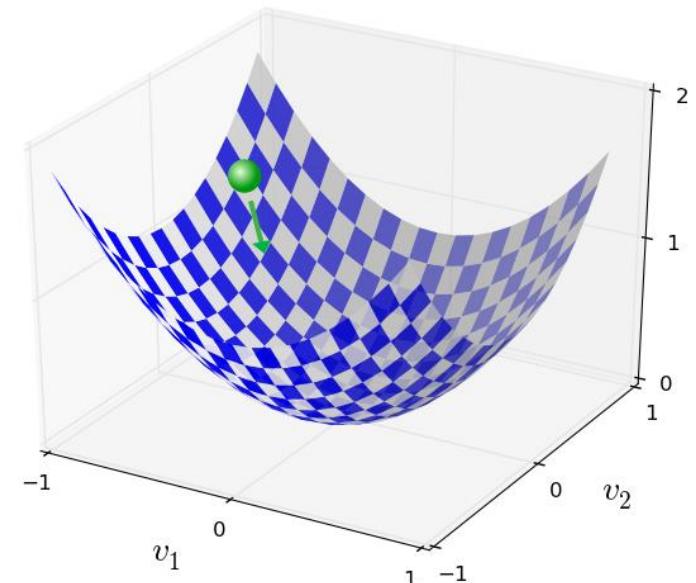


Ground truth for “6”:

$$y(x) = (0, 0, 0, 0, 0, 0, 1, 0, 0, 0)^T$$

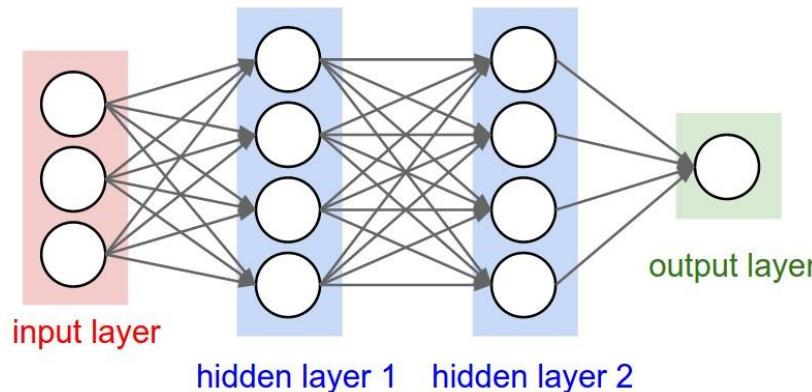
“Loss” function:

$$C(w, b) \equiv \frac{1}{2n} \sum_x \|y(x) - a\|^2$$

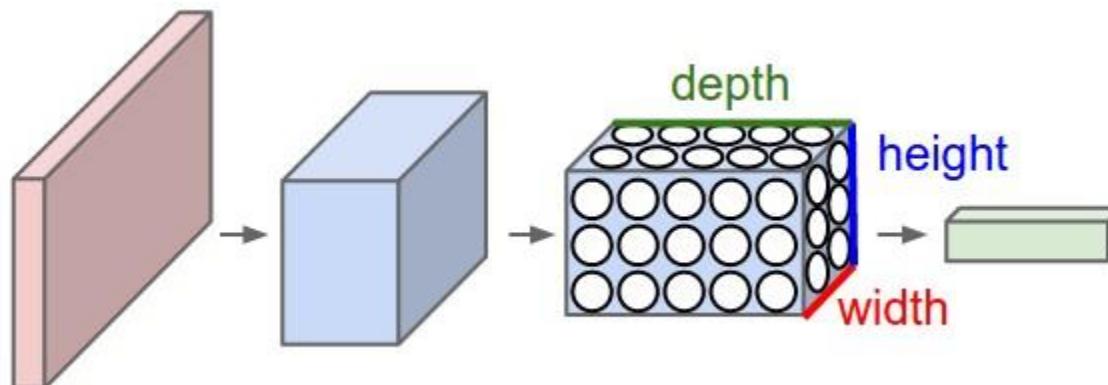


# Convolutional Neural Networks

Regular neural network (fully connected):



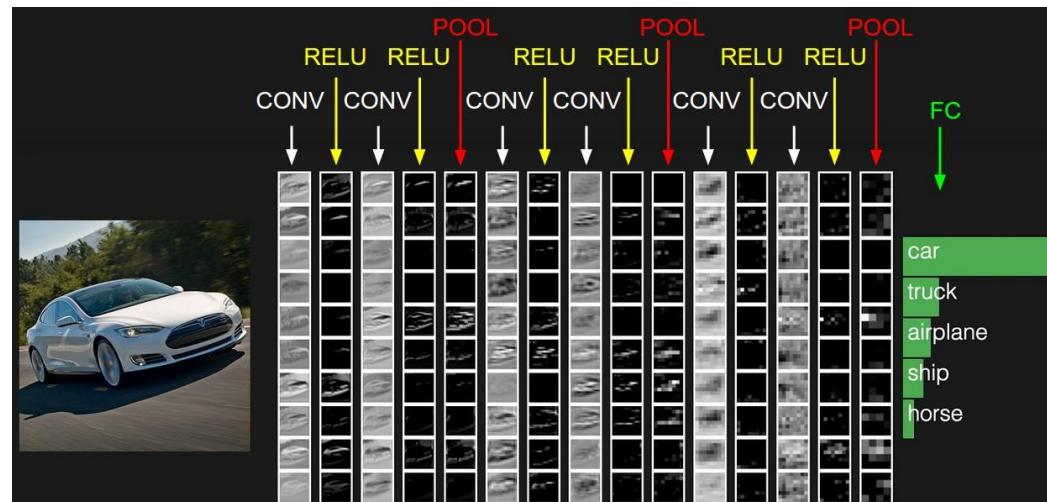
Convolutional neural network:



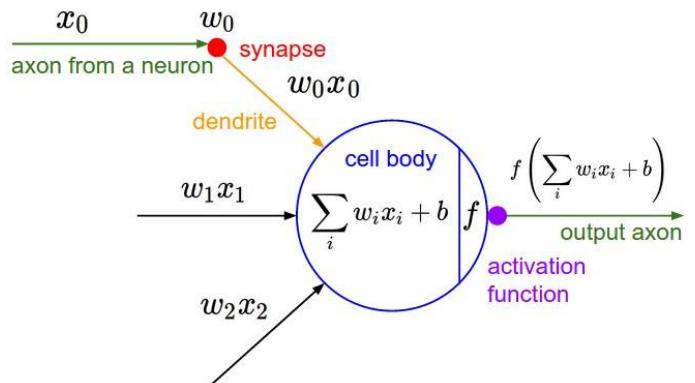
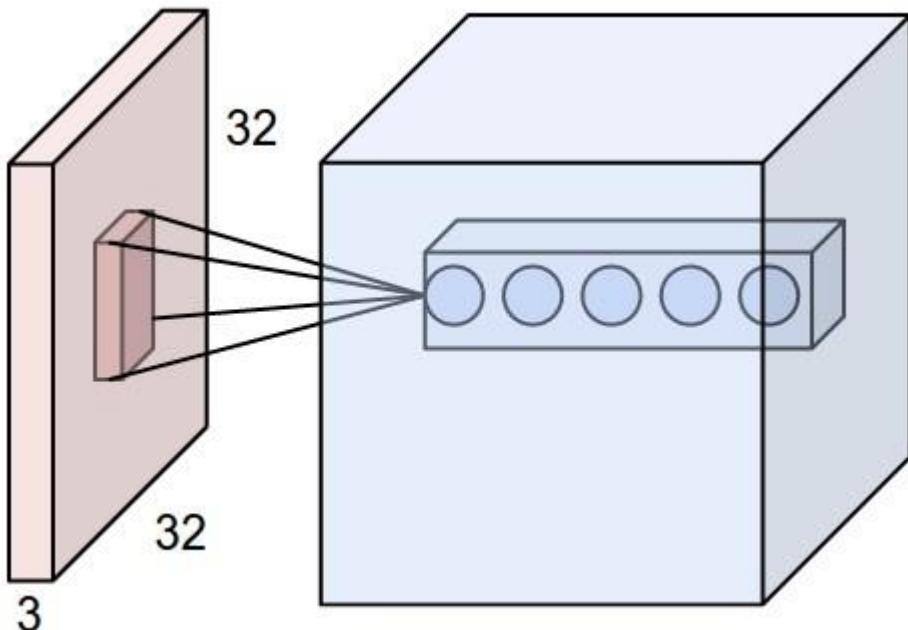
Each layer takes a 3d volume, produces 3d volume with some smooth function that may or may not have parameters.

# Convolutional Neural Networks: Layers

- **INPUT** [32x32x3] will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R,G,B.
- **CONV** layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as [32x32x12] if we decided to use 12 filters.
- **RELU** layer will apply an elementwise activation function, such as the  $\max(0,x)$  thresholding at zero. This leaves the size of the volume unchanged ([32x32x12]).
- **POOL** layer will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as [16x16x12].
- **FC** (i.e. fully-connected) layer will compute the class scores, resulting in volume of size [1x1x10], where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.



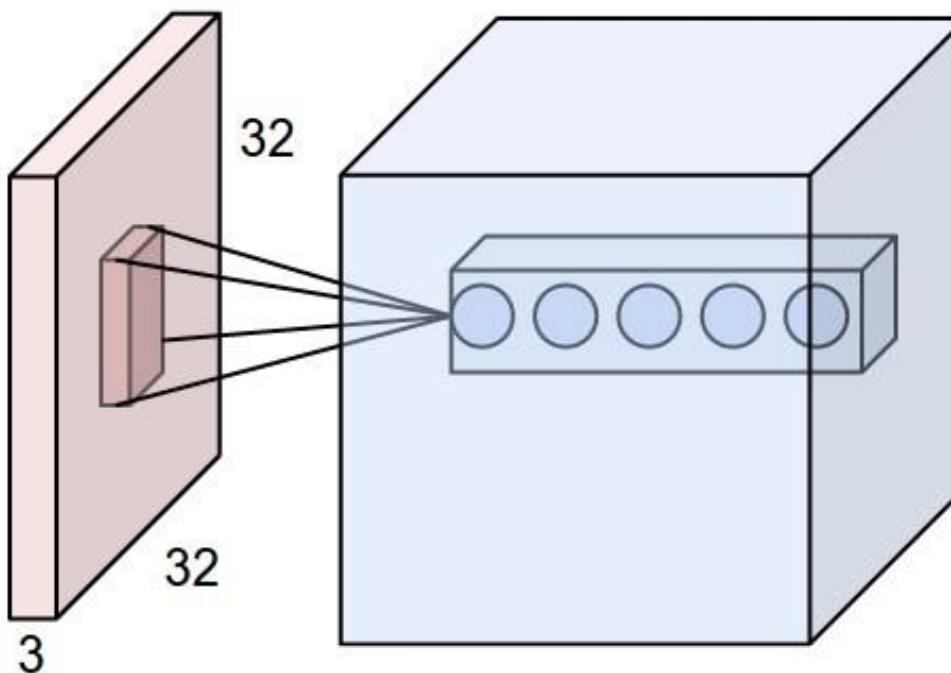
# Dealing with Images: Local Connectivity



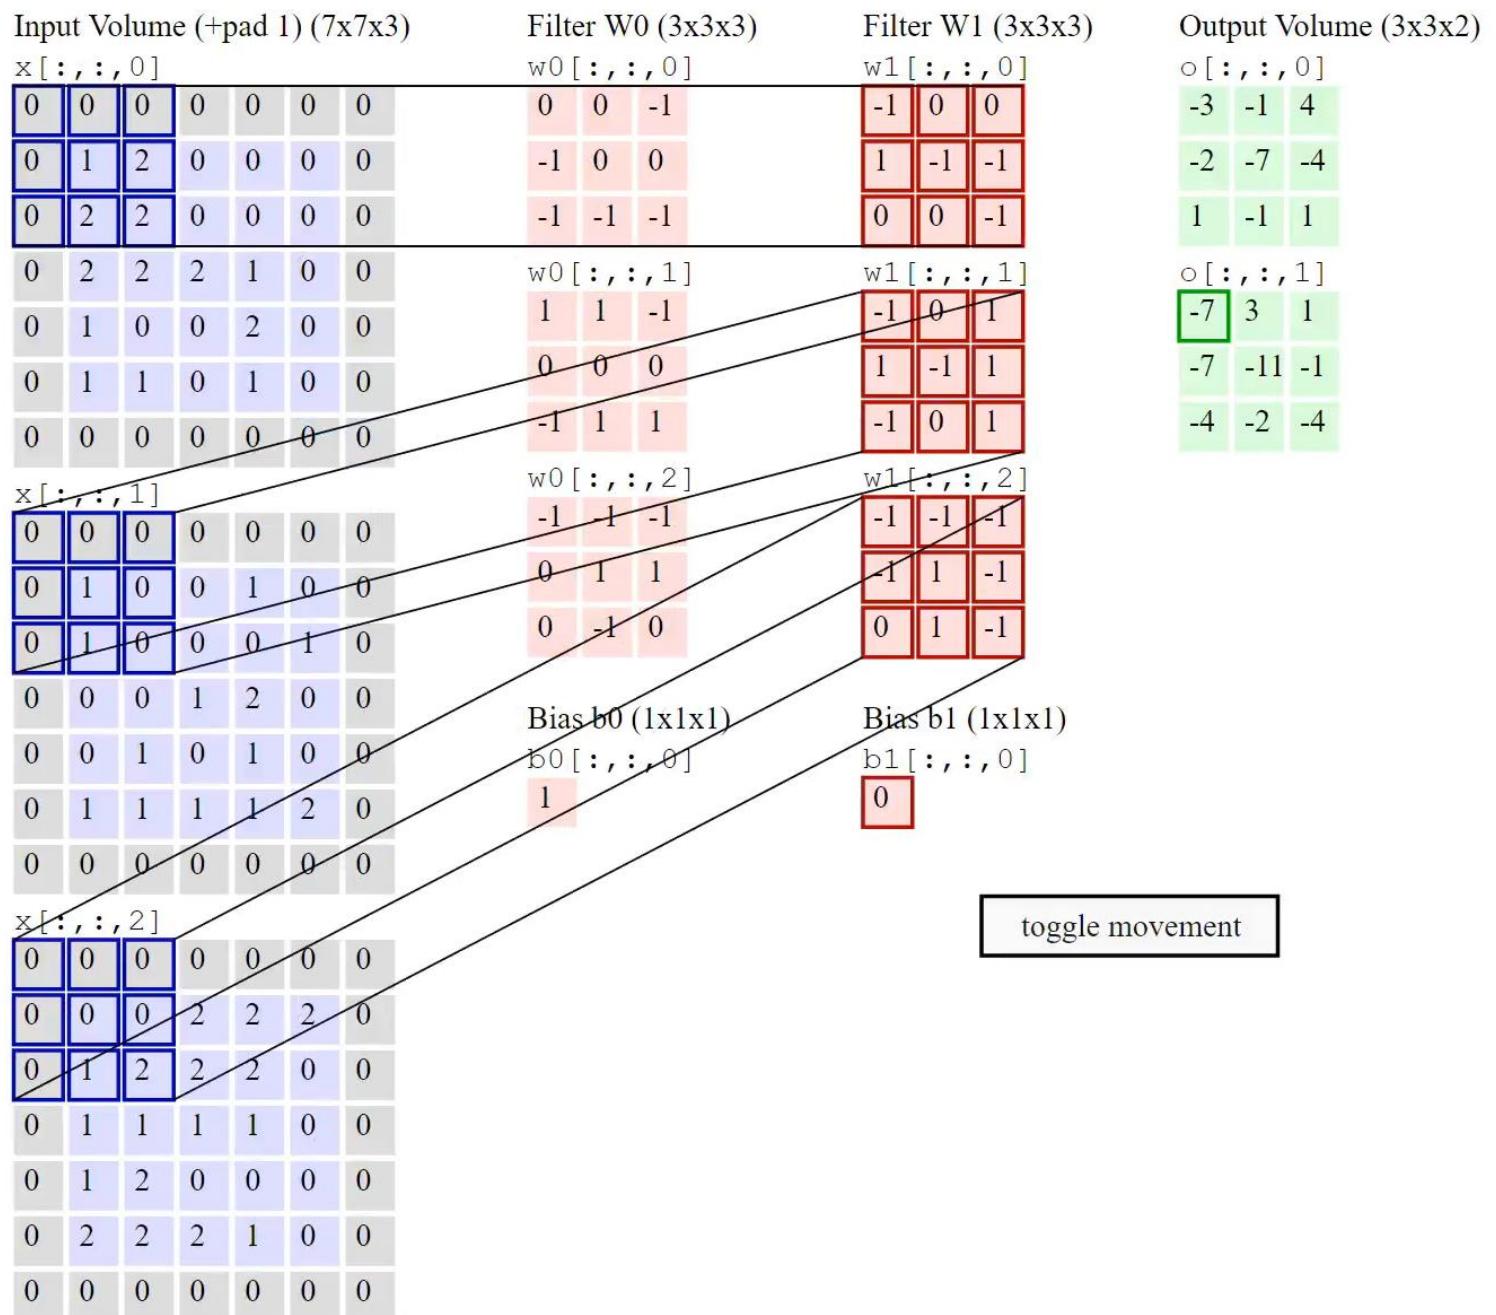
Same neuron. Just more focused (narrow “receptive field”).

The parameters on each filter are spatially “shared”  
(if a feature is useful in one place, it’s useful elsewhere)

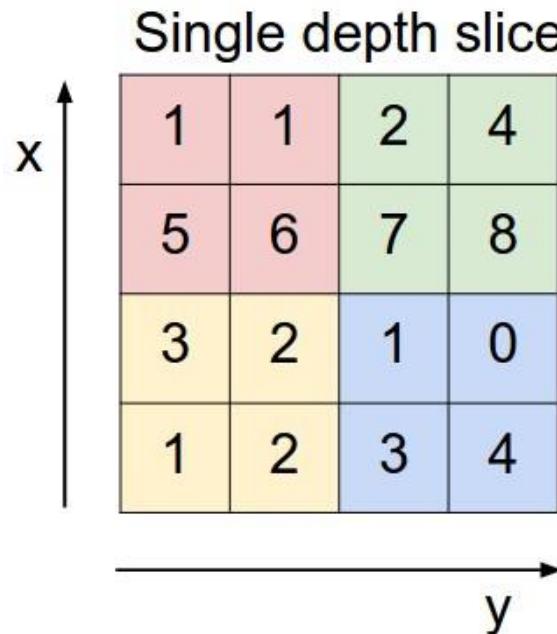
# ConvNets: Spatial Arrangement of Output Volume



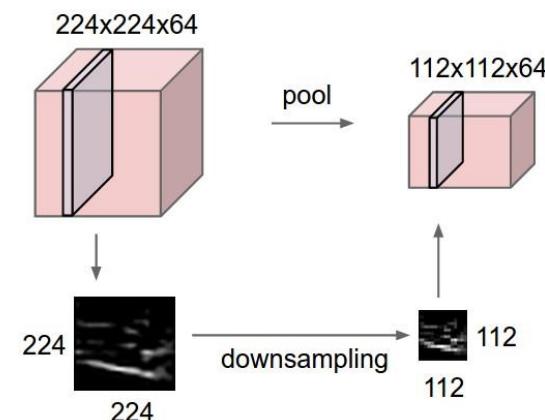
- **Depth:** number of filters
- **Stride:** filter step size (when we “slide” it)
- **Padding:** zero-pad the input



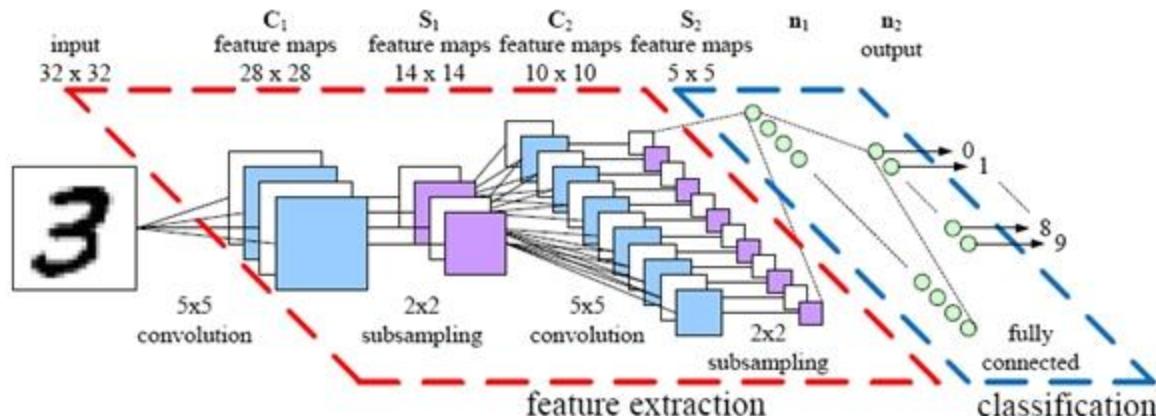
# ConvNets: Pooling



max pool with 2x2 filters  
and stride 2

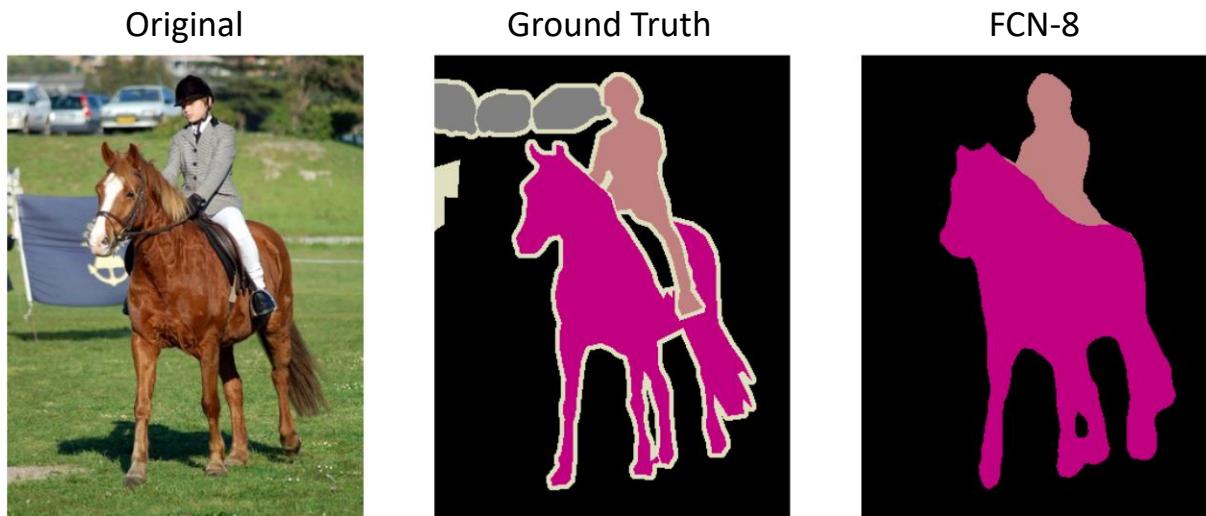
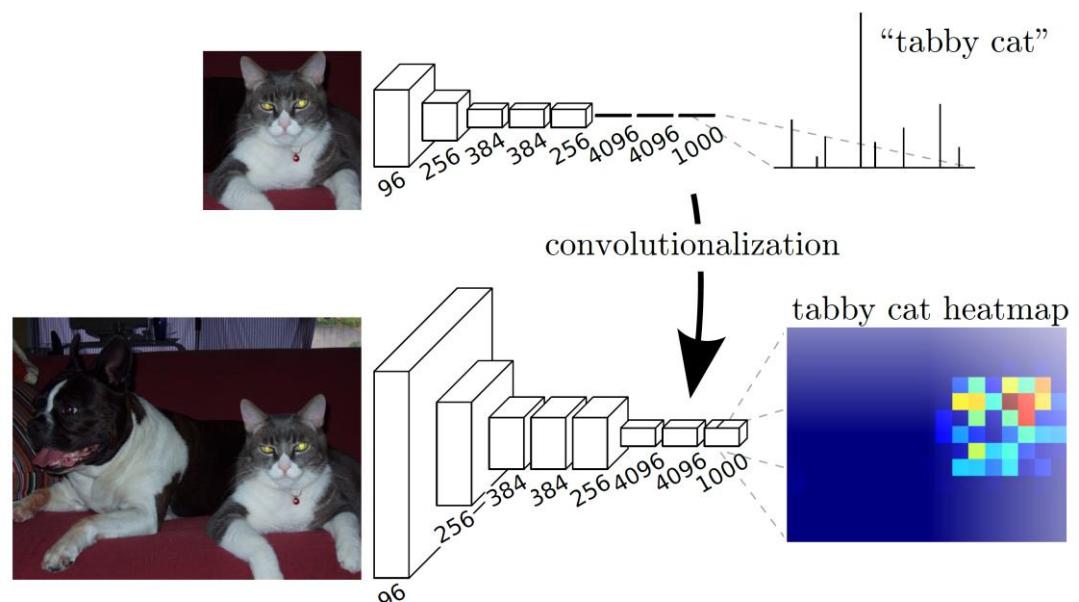


# Computer Vision: Object Recognition / Classification

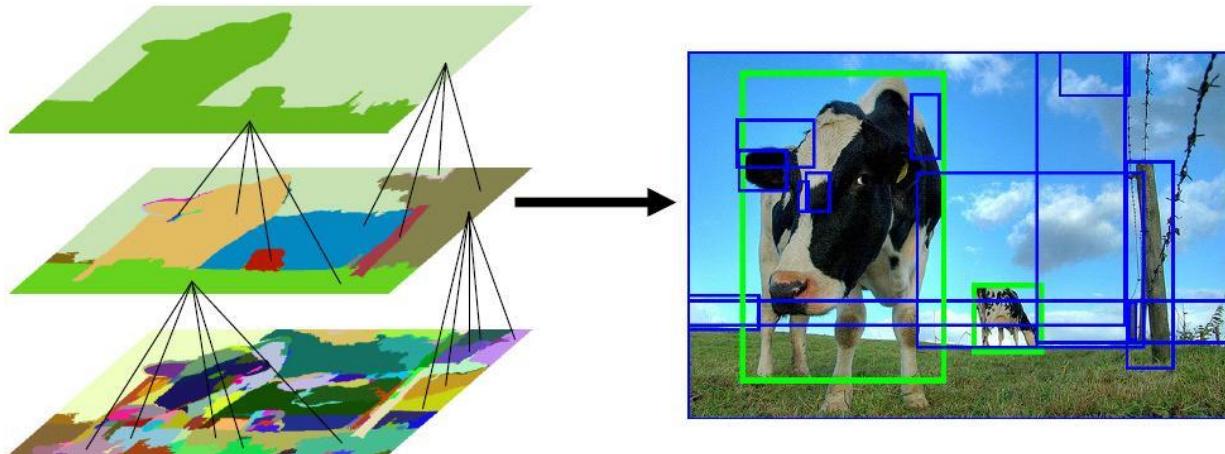


mite	container ship	motor scooter	leopard
mite	container ship	motor scooter	leopard
black widow	lifeboat	go-kart	jaguar
cockroach	amphibian	moped	cheetah
tick	fireboat	bumper car	snow leopard
starfish	drilling platform	golfcart	Egyptian cat

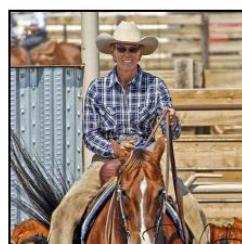
# Computer Vision: Segmentation



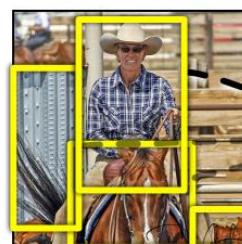
# Computer Vision: Object Detection



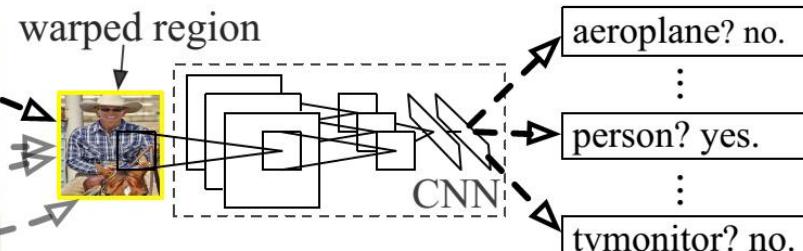
## R-CNN: *Regions with CNN features*



1. Input  
image



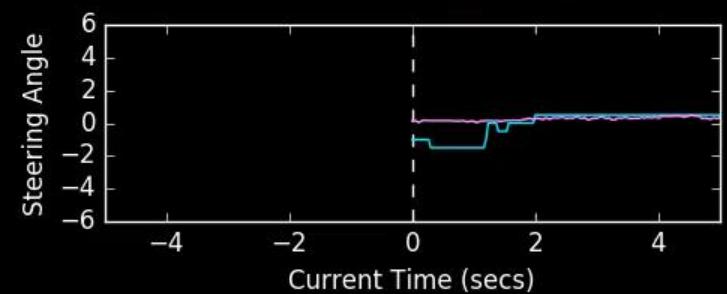
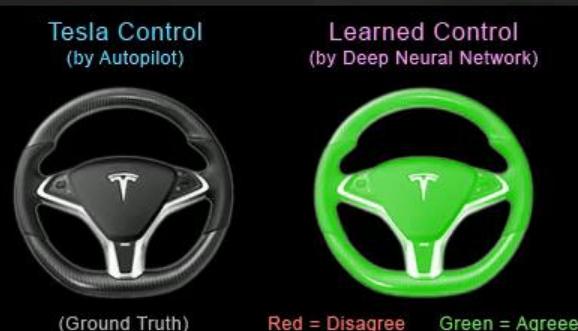
2. Extract region  
proposals (~2k)



3. Compute  
CNN features

4. Classify  
regions

# How Can Convolutional Neural Networks Help Us Drive?



# Driving: The Numbers

(in United States, in 2014)

## Miles

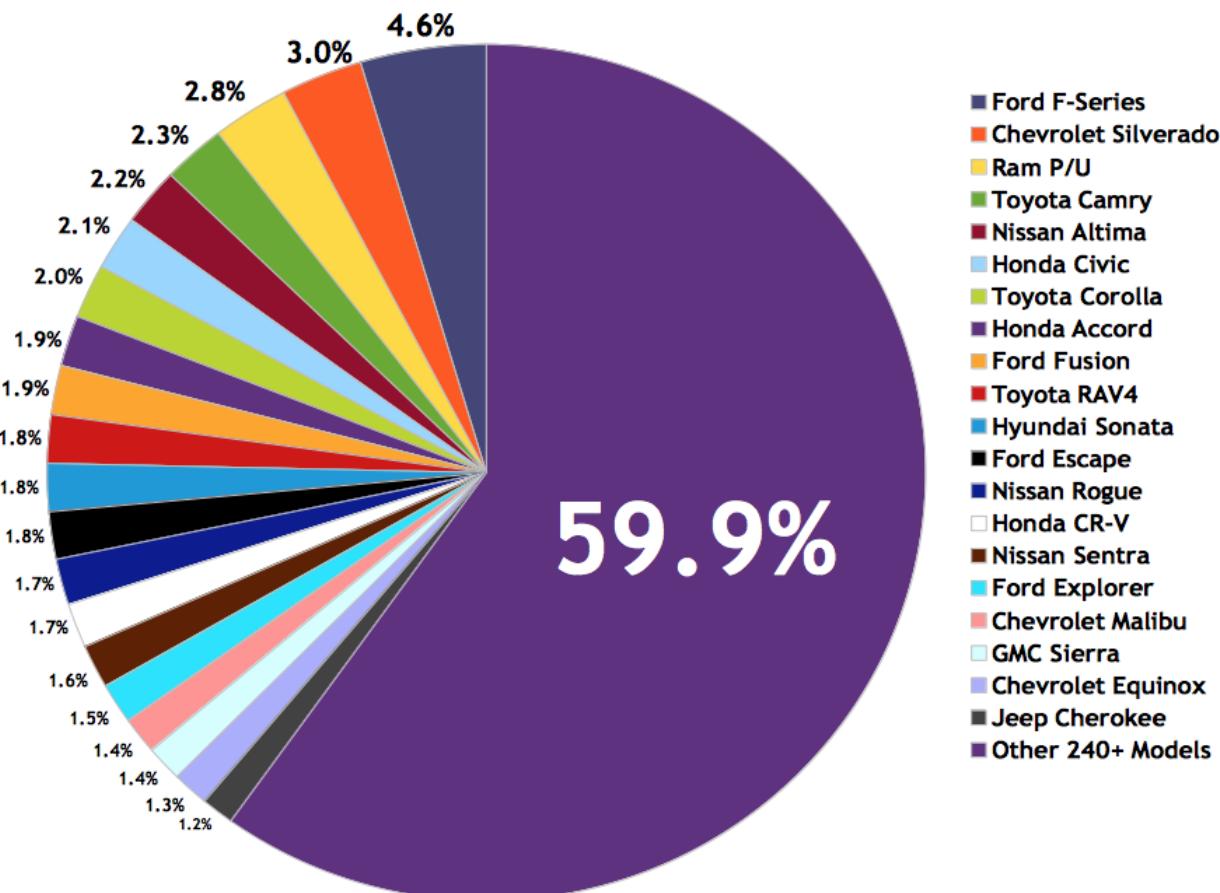
- **All drivers:** 10,658 miles  
(29.2 miles per day)
- **Rural drivers:** 12,264 miles
- **Urban drivers:** 9,709 miles

## Fatalities

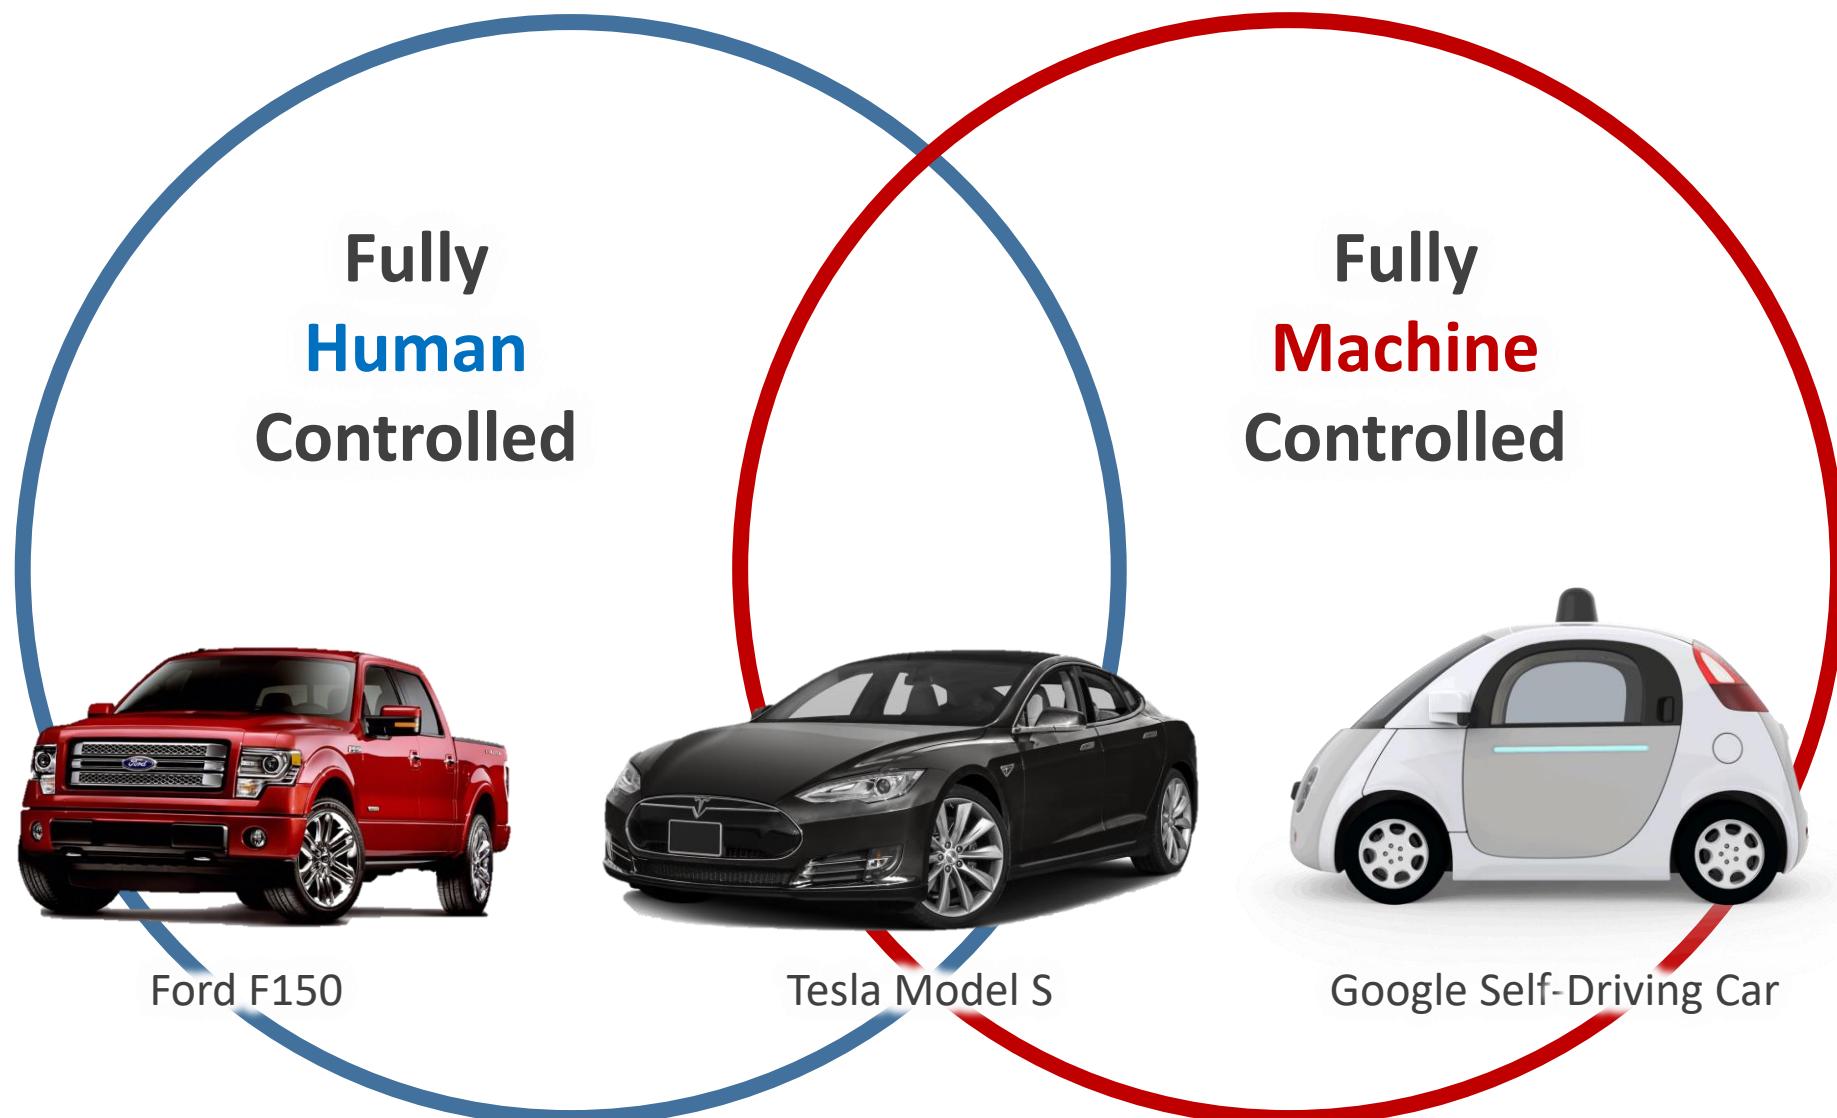
- **Fatal crashes:** 29,989
- **All fatalities:** 32,675
- **Car occupants:** 12,507
- **SUV occupants:** 8,320
- **Pedestrians:** 4,884
- **Motorcycle:** 4,295
- **Bicyclists:** 720
- **Large trucks:** 587

# Cars We Drive

Market Share Of America's 20 Best-Selling Vehicles  
March 2016

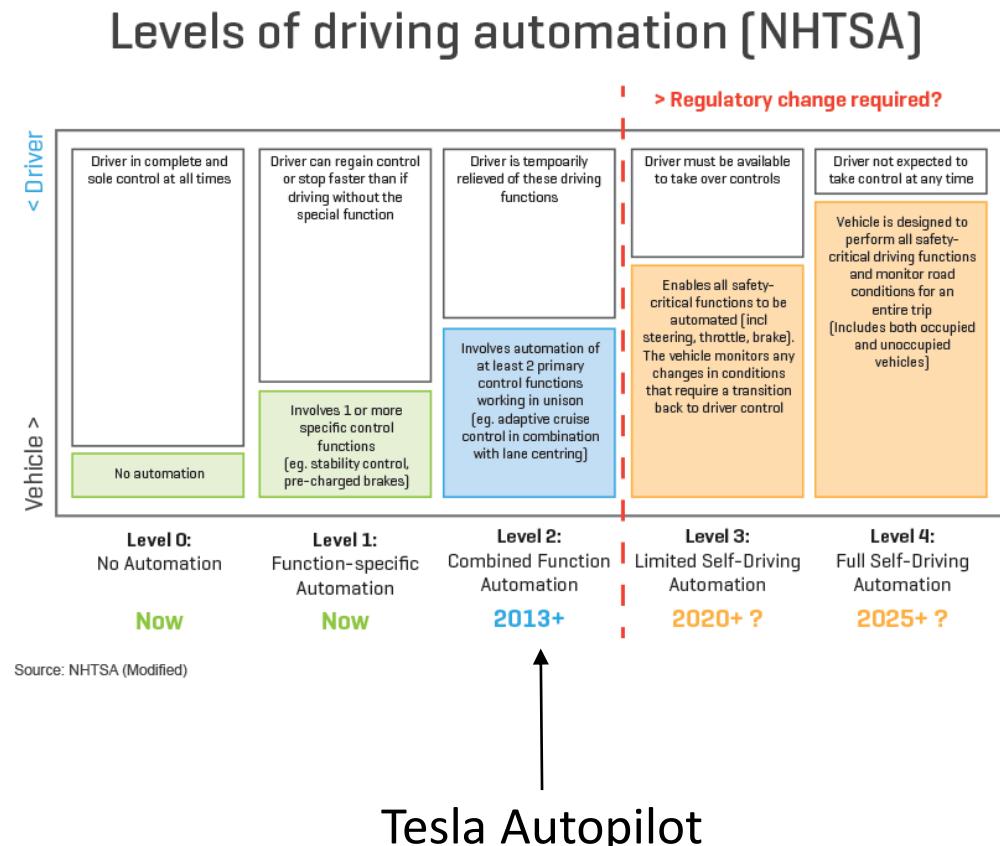


# **Human** at the Center of Automation: The Way to Full Autonomy Includes the Human



# Human at the Center of Automation: The Way to Full Autonomy Includes the Human

- **Emergency**
  - Automatic emergency breaking (AEB)
- **Warnings**
  - Lane departure warning (LDW)
  - Forward collision warning (FCW)
  - Blind spot detection
- **Longitudinal**
  - Adaptive cruise control (ACC)
- **Lateral**
  - Lane keep assist (LKA)
  - Automatic steering
- **Control and Planning**
  - Automatic lane change
  - Automatic parking



# Distracted Humans

## What is distracted driving?

- Texting
- Using a smartphone
- Eating and drinking
- Talking to passengers
- Grooming
- Reading, including maps
- Using a navigation system
- Watching a video
- Adjusting a radio

- **Injuries and fatalities:**

3,179 people were killed and 431,000 were injured in motor vehicle crashes involving distracted drivers  
*(in 2014)*

- **Texts:**

169.3 billion text messages were sent in the US every month.  
*(as of December 2014)*

- **Eye off road:**

5 seconds is the average time your eyes are off the road while texting. When traveling at 55mph, that's enough time to cover the length of a football field blindfolded.

# 4 D's of Being Human: Drunk, Drugged, Distracted, Drowsy

- **Drunk Driving:** In 2014, 31 percent of traffic fatalities involved a drunk driver.
- **Drugged Driving:** 23% of night-time drivers tested positive for illegal, prescription or over-the-counter medications.
- **Distracted Driving:** In 2014, 3,179 people (10 percent of overall traffic fatalities) were killed in crashes involving distracted drivers.
- **Drowsy Driving:** In 2014, nearly three percent of all traffic fatalities involved a drowsy driver, and at least 846 people were killed in crashes involving a drowsy driver.

# *In Context:* Traffic Fatalities

Total miles driven in U.S. in 2014:

3,000,000,000,000 (3 million million)

Tesla Autopilot miles driven since October 2015:

300,000,000 (300 million)

(as of December 2016)

# *In Context: Traffic Fatalities*

Total miles driven in U.S. in 2014:

3,000,000,000,000 (3 million million)

Fatalities: **32,675**

(1 in 90 million)

Tesla Autopilot miles driven since October 2015:

300,000,000 (300 million)

Fatalities: **1**

# *In Context: Traffic Fatalities*

We (increasingly) understand this → **Fatalities: 32,675**

Total miles driven in U.S. in 2014:  
3,000,000,000,000 (3 million million)

(1 in 90 million)

We do not understand this (yet) → **Fatalities: 1**

Tesla Autopilot miles driven since October 2015:

300,000,000 (300 million)

**We need A LOT of real-world semi-autonomous driving data!**

Computer Vision + Machine Learning + Big Data = Understanding

# The Data



Teslas instrumented: 17

Hours of data: 5,000+ hours

Distance traveled: 70,000+ miles



# The Data



Total Time Driving: 0 mins  
Autopilot Available: 0 mins  
Autopilot Engaged: 0 mins



# Camera and Lens Selection



**Logitech C920:**  
On-board H264 Compression



**Fisheye:** Capture full range of head, body movement inside vehicle.

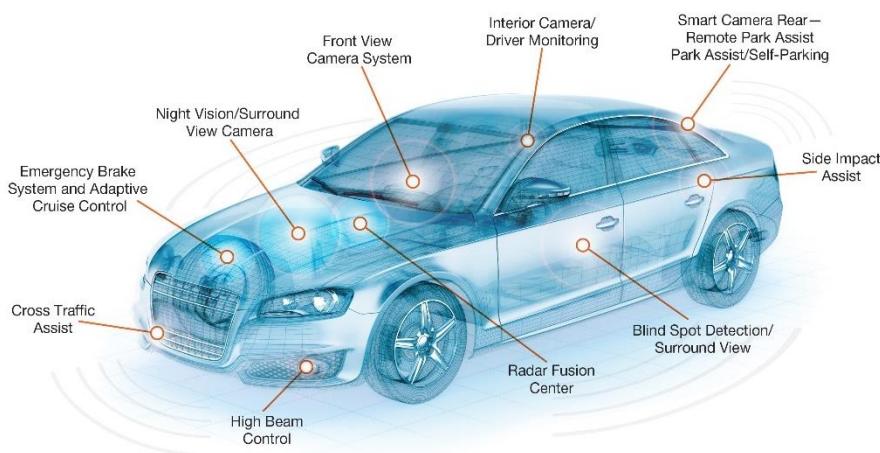
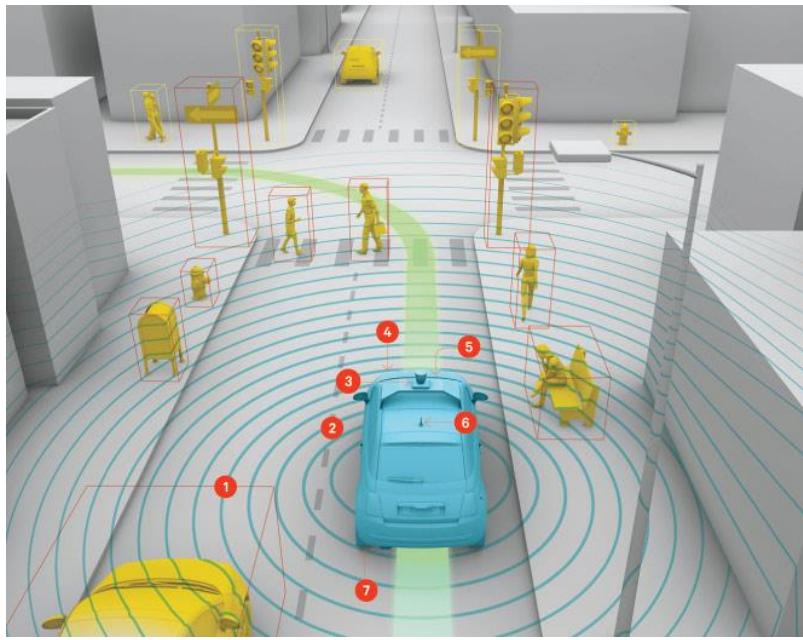


**Case for C-Mount Lens:**  
Flexibility in lens selection



**2.8-12mm Focal Length:** “Zoom” on the face without obstructing the driver’s view.

# Semi-Autonomous Vehicle Components



## External

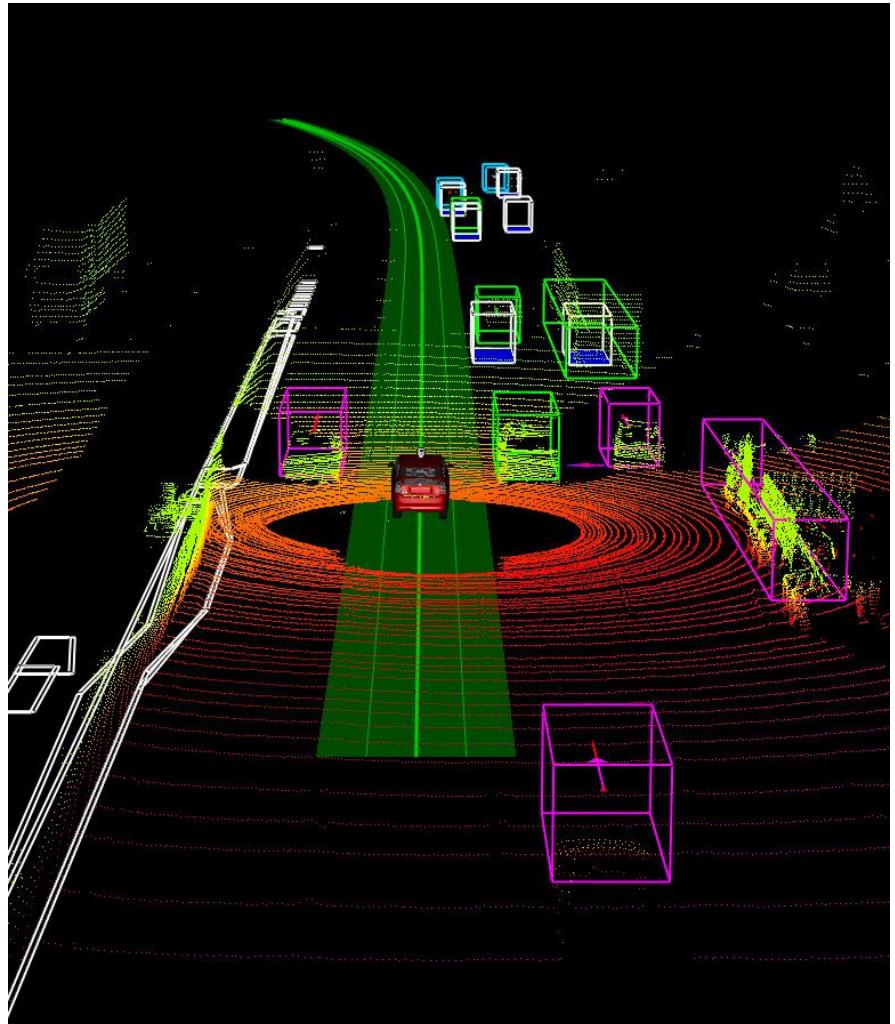
1. Radar
2. Visible-light camera
3. LIDAR
4. Infrared camera
5. Stereo vision
6. GPS/IMU
7. CAN
8. Audio

## Internal

1. Visible-light camera
2. Infrared camera
3. Audio

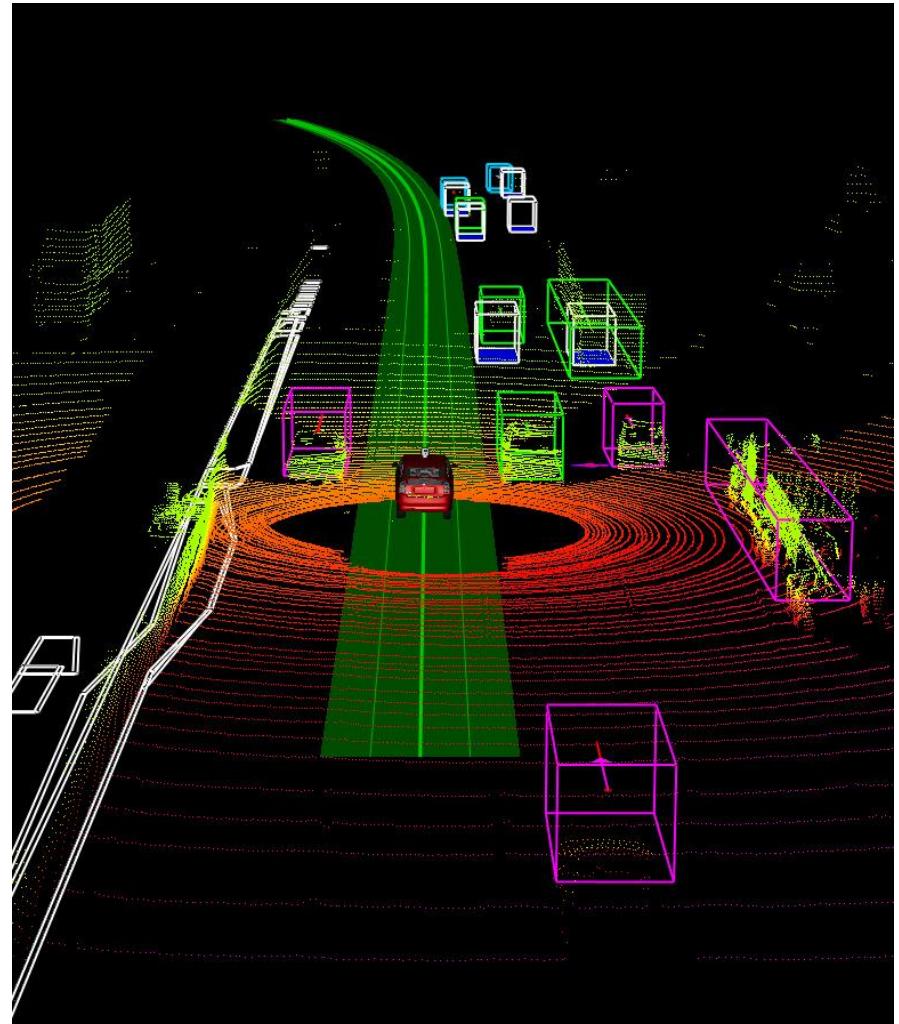
# Self-Driving Car Tasks

- **Localization and Mapping:**  
Where am I?
- **Scene Understanding:**  
Where is everyone else?
- **Movement Planning:**  
How do I get from A to B?
- **Driver State:**  
What's the driver up to?



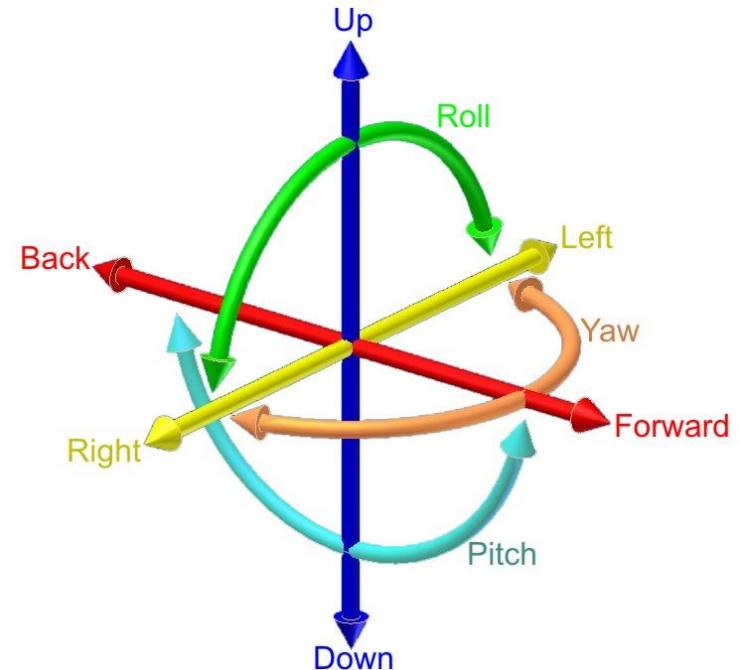
# Self-Driving Car Tasks

- **Localization and Mapping:**  
Where am I?
- **Scene Understanding:**  
Where is everyone else?
- **Movement Planning:**  
How do I get from A to B?
- **Driver State:**  
What's the driver up to?



# Visual Odometry

- 6-DOF: freedom of movement
  - Changes in position:
    - Forward/backward: surge
    - Left/right: sway
    - Up/down: heave
  - Orientation:
    - Pitch, Yaw, Roll
- Source:
  - **Monocular:** I moved 1 unit
  - **Stereo:** I moved 1 meter
  - Mono = Stereo for far away objects
    - PS: For tiny robots everything is “far away” relative to inter-camera distance



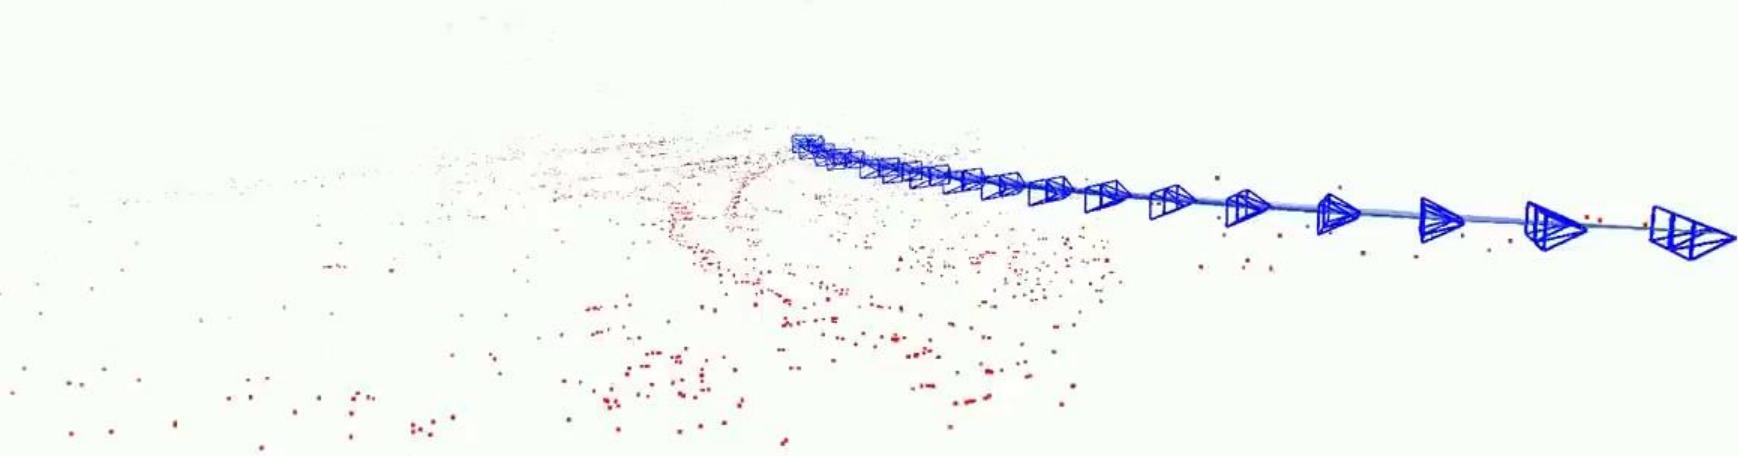
# SLAM: Simultaneous Localization and Mapping

*What works: SIFT and optical flow*

x2



TRACKING — KFs: 19 , MPs: 3527 , Tracked: 369 + 0



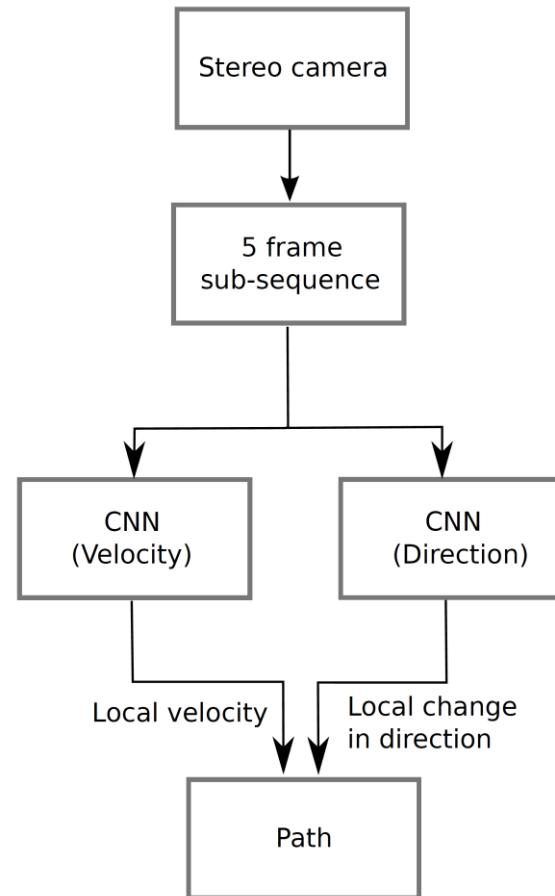
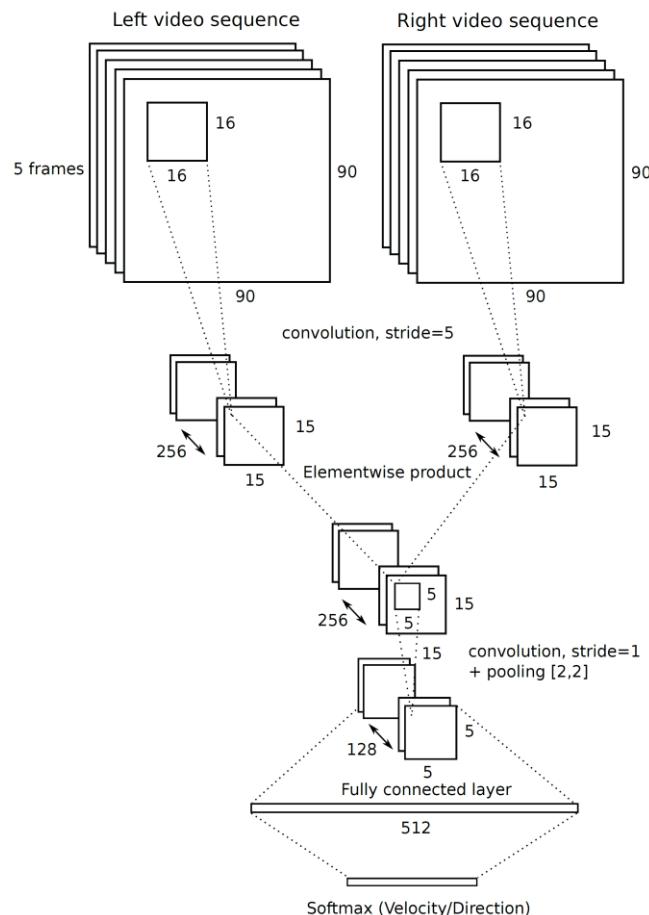
# Visual Odometry in Parts



- (Stereo) Undistortion, Rectification
- (Stereo) Disparity Map Computation
- Feature Detection (e.g., SIFT, FAST)
- Feature Tracking (e.g., KLT: Kanade-Lucas-Tomasi)
- Trajectory Estimation
  - Use rigid parts of the scene (requires outlier/inlier detection)
  - For mono, need more info\* like camera orientation and height off the ground

\* Kitt, Bernd Manfred, et al. "Monocular visual odometry using a planar road model to solve scale ambiguity." (2011).

# End-to-End Visual Odometry



Konda, Kishore, and Roland Memisevic. "Learning visual odometry with a convolutional network." *International Conference on Computer Vision Theory and Applications*. 2015.

# Self-Driving Car Tasks

- **Localization and Mapping:**

Where am I?

- **Scene Understanding:**

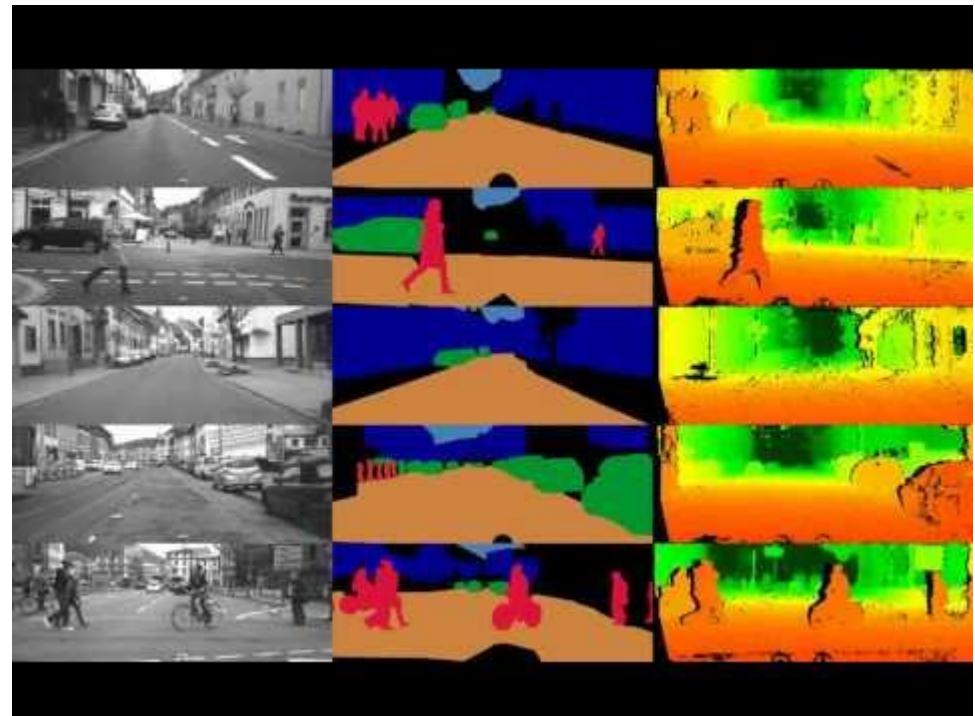
Where is everyone else?

- **Movement Planning:**

How do I get from A to B?

- **Driver State:**

What's the driver up to?



# Object Detection



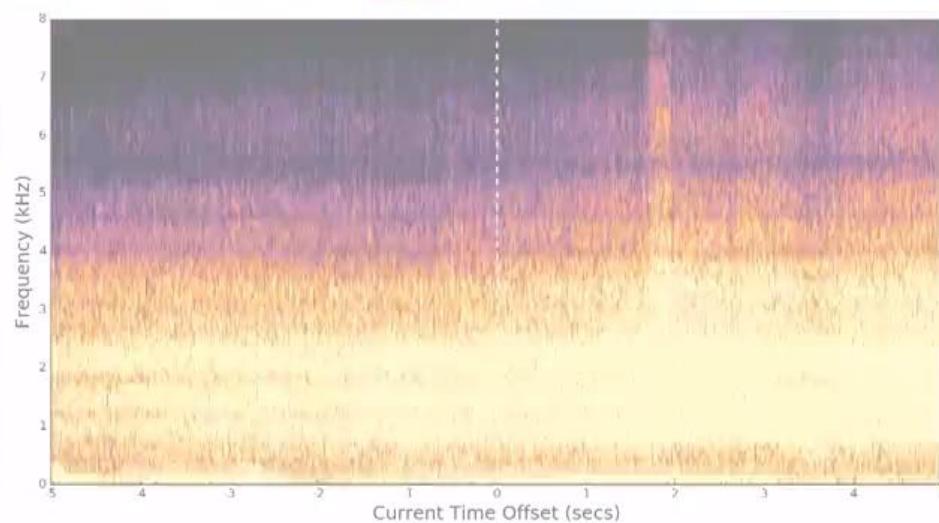
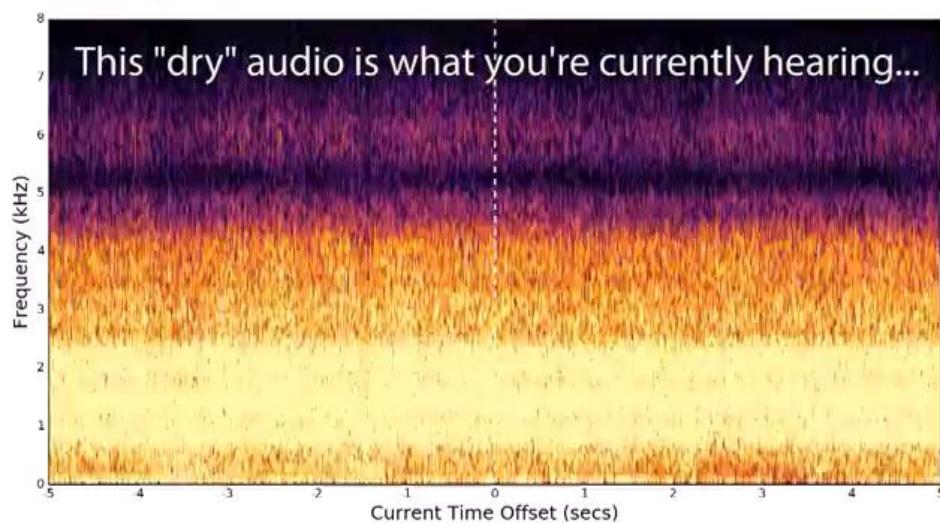
- Past approaches: cascades classifiers (Haar-like features)
- Where deep learning can help:  
recognition, classification, detection

# Full Driving Scene Segmentation



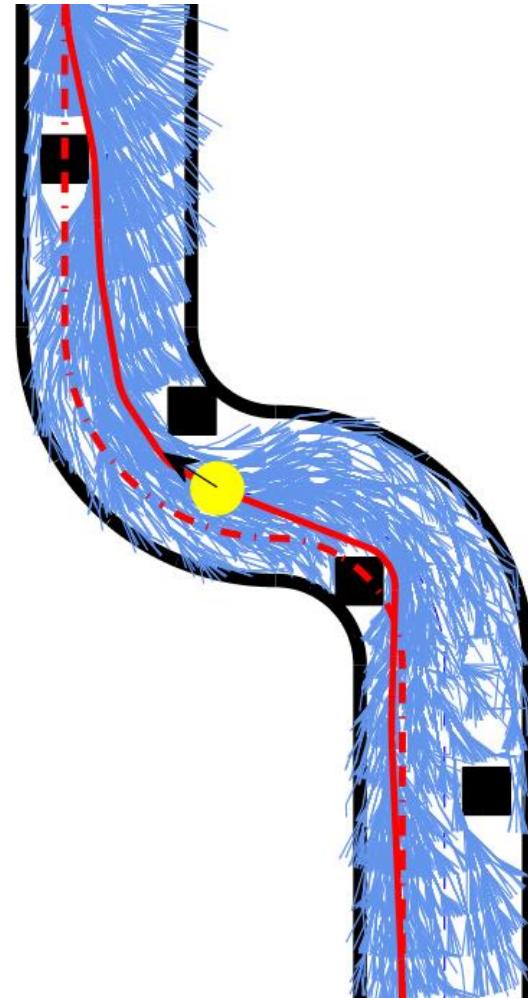
Fully Convolutional Network implementation:  
<https://github.com/tkuanlun350/Tensorflow-SegNet>

# Road Texture and Condition from Audio (with Recurrent Neural Networks)

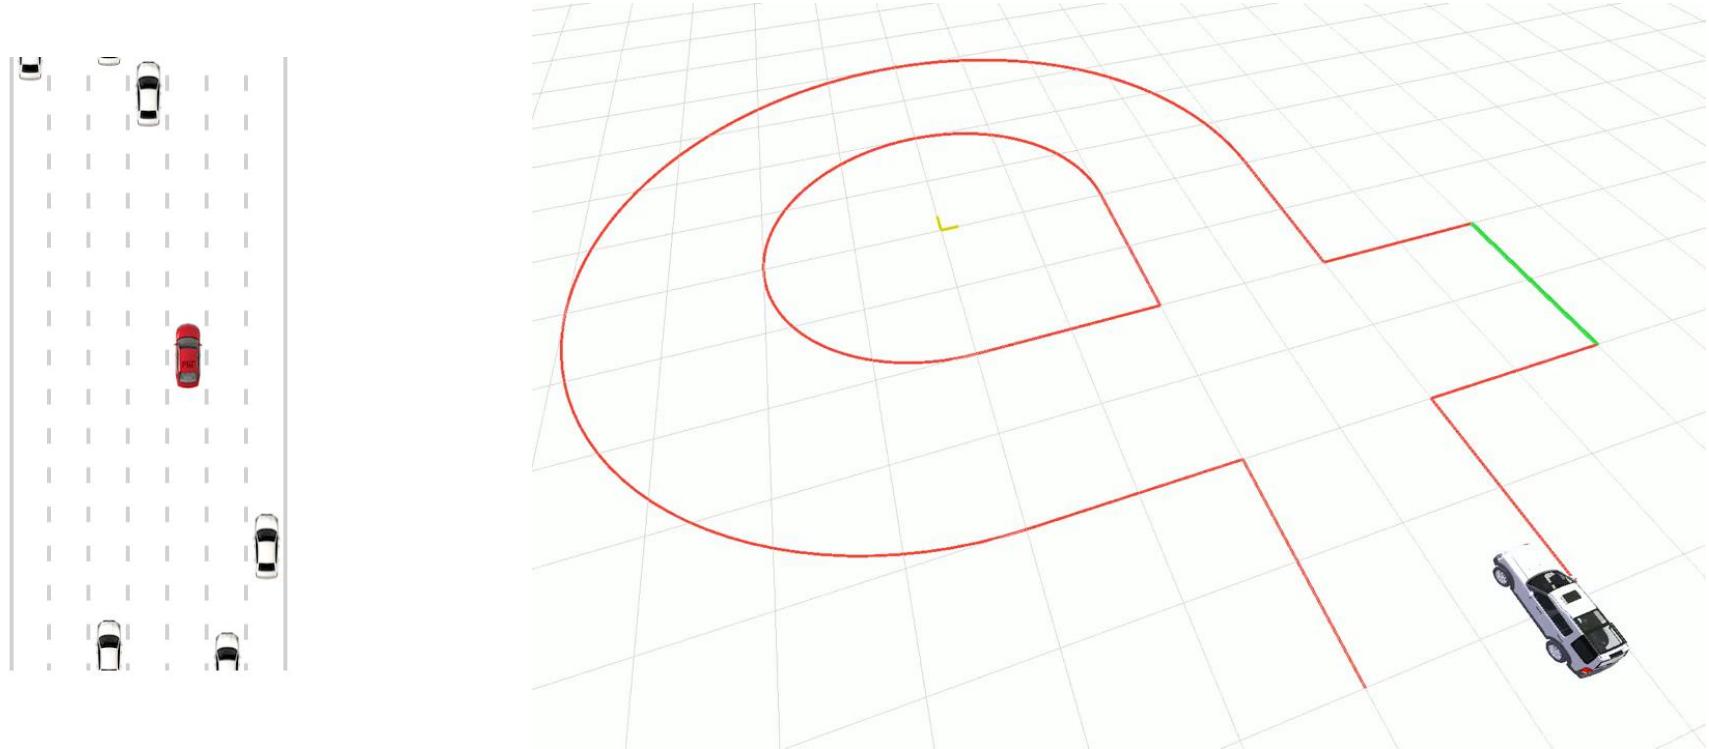


# Self-Driving Car Tasks

- **Localization and Mapping:**  
Where am I?
- **Scene Understanding:**  
Where is everyone else?
- **Movement Planning:**  
How do I get from A to B?
- **Driver State:**  
What's the driver up to?



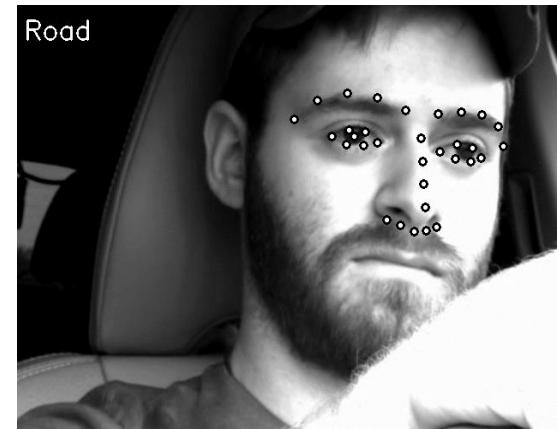
- Previous approaches: optimization-based control
- Where deep learning can help: reinforcement learning



Deep Reinforcement Learning implementation:  
<https://github.com/nivwusquorum/tensorflow-deepq>

# Self-Driving Car Tasks

- **Localization:**  
Where am I?
- **Object detection:**  
Where is everyone else?
- **Movement planning:**  
How do I get from A to B?
- **Driver state:**  
What's the driver up to?



# Drive State Detection: A Multi-Resolutonal View

Increasing level of detection resolution and difficulty



Body Pose

Head Pose

Blink Rate

Blink Duration

Eye Pose

Blink Dynamics

Pupil Diameter

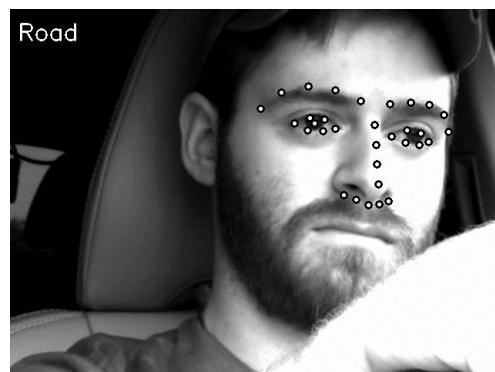
Micro Saccades

Gaze Classification

Drowsiness

Micro Glances

Cognitive Load



Frames: 1

Time: 0.03 secs

Total Confident Decisions: 1

Correct Confident Decisions: 1

Wrong Confident Decisions: 0

Accuracy: 100%



Frames: 1

Time: 0.03 secs

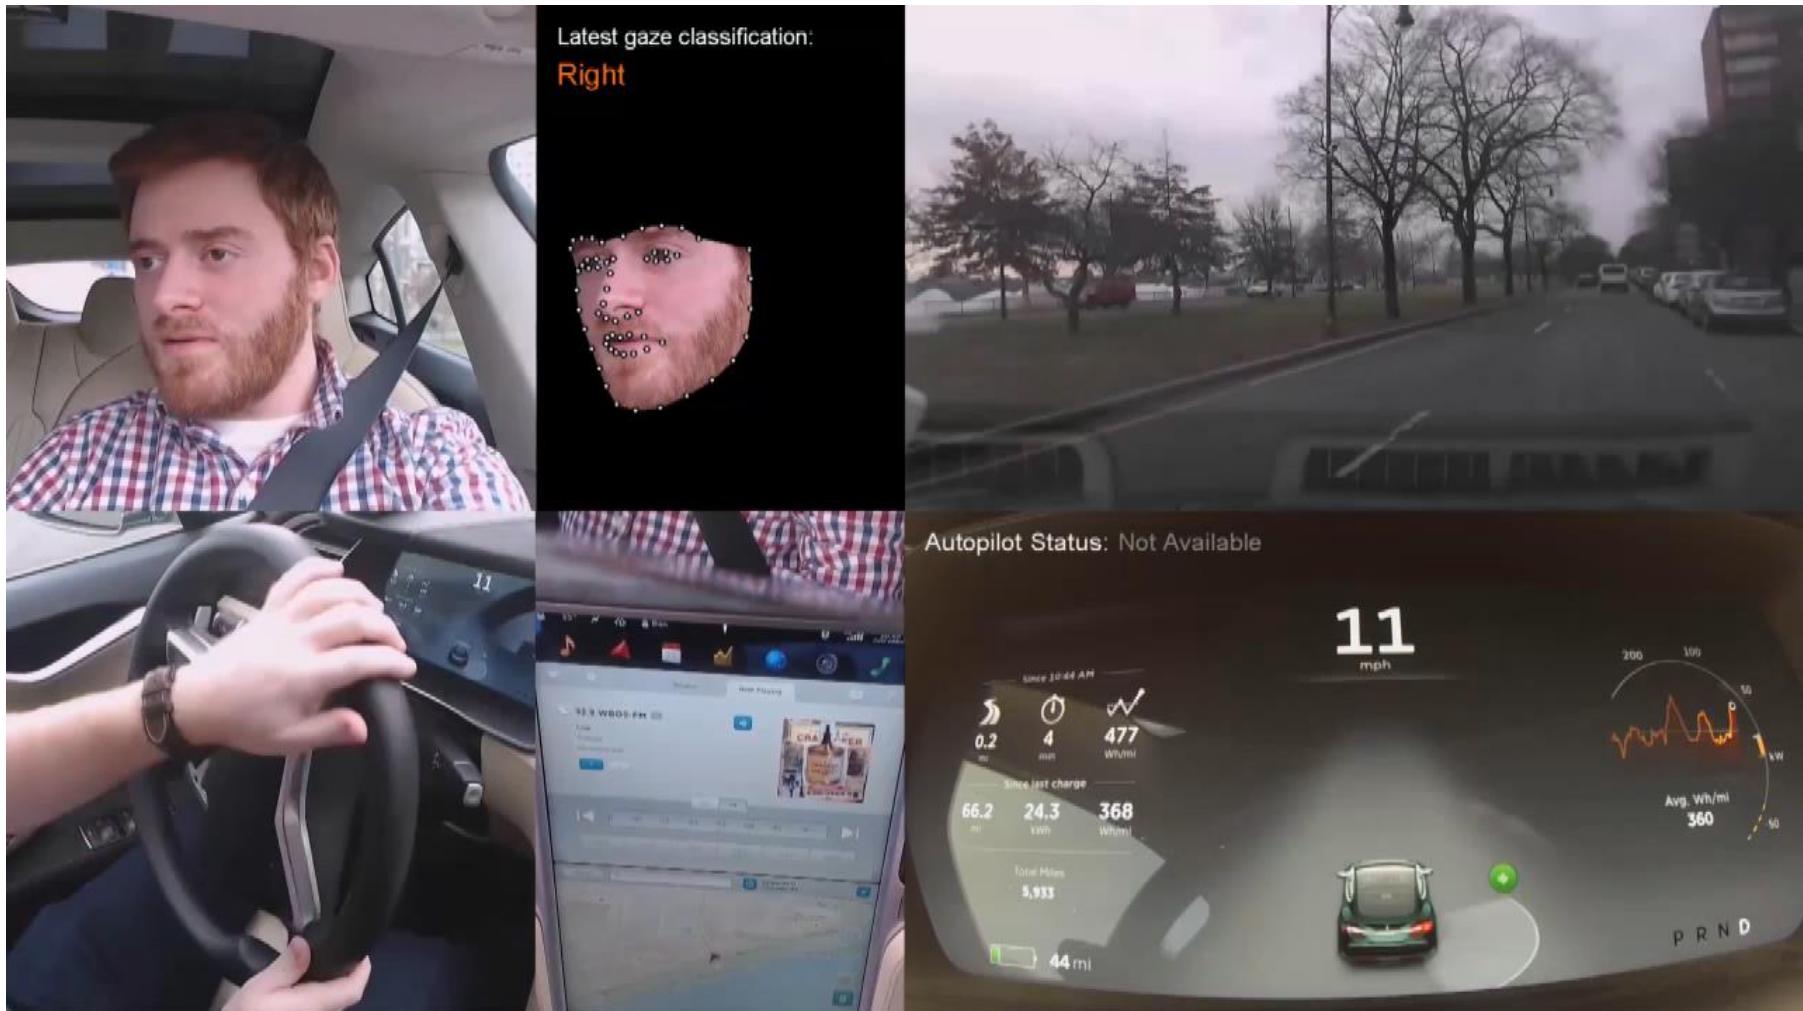
Total Confident Decisions: 1

Correct Confident Decisions: 1

Wrong Confident Decisions: 0

Accuracy: 100%

# Gaze Region and Autopilot State



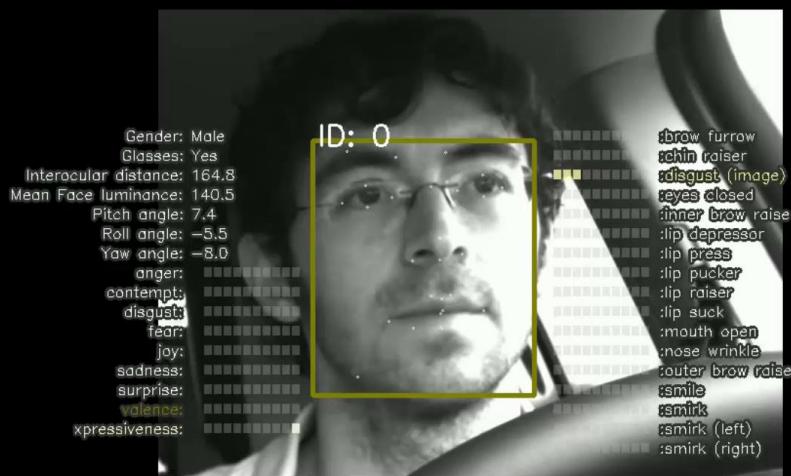
# Driver Emotion

Class 1: **Satisfied** with Voice-Based Interaction



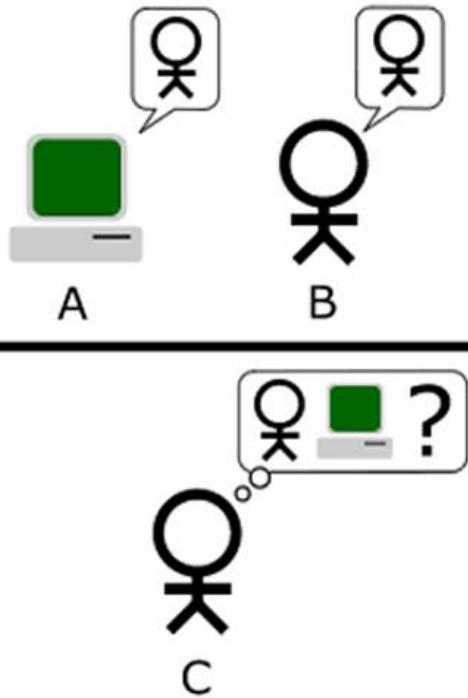
:brow furrow  
:chin raiser  
:disgust (image)  
:eyes closed  
:inner brow raise  
:lip depressor  
:lip press  
:lip pucker  
:lip raiser  
:lip suck  
:mouth open  
:nose wrinkle  
:outer brow raise  
:smile  
:smirk  
:smirk (left)  
:smirk (right)

Class 2: **Frustrated** with Voice-Based Interaction



:brow furrow  
:chin raiser  
:disgust (image)  
:eyes closed  
:inner brow raise  
:lip depressor  
:lip press  
:lip pucker  
:lip raiser  
:lip suck  
:mouth open  
:nose wrinkle  
:outer brow raise  
:smile  
:smirk  
:smirk (left)  
:smirk (right)

If Driving is a Conversation, this is an  
End-to-End Natural Language Generation



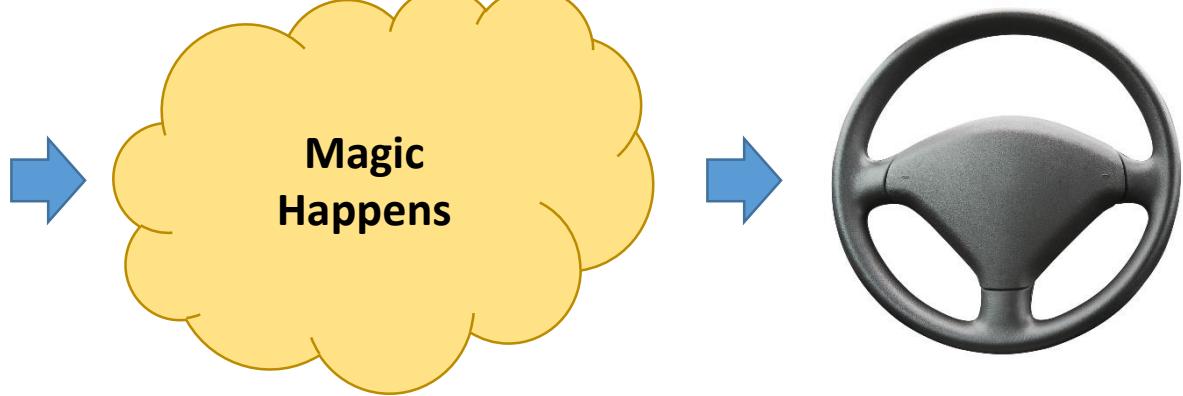
**Turing Test:**

Can a computer be mistaken for a human more than 30% of the time?

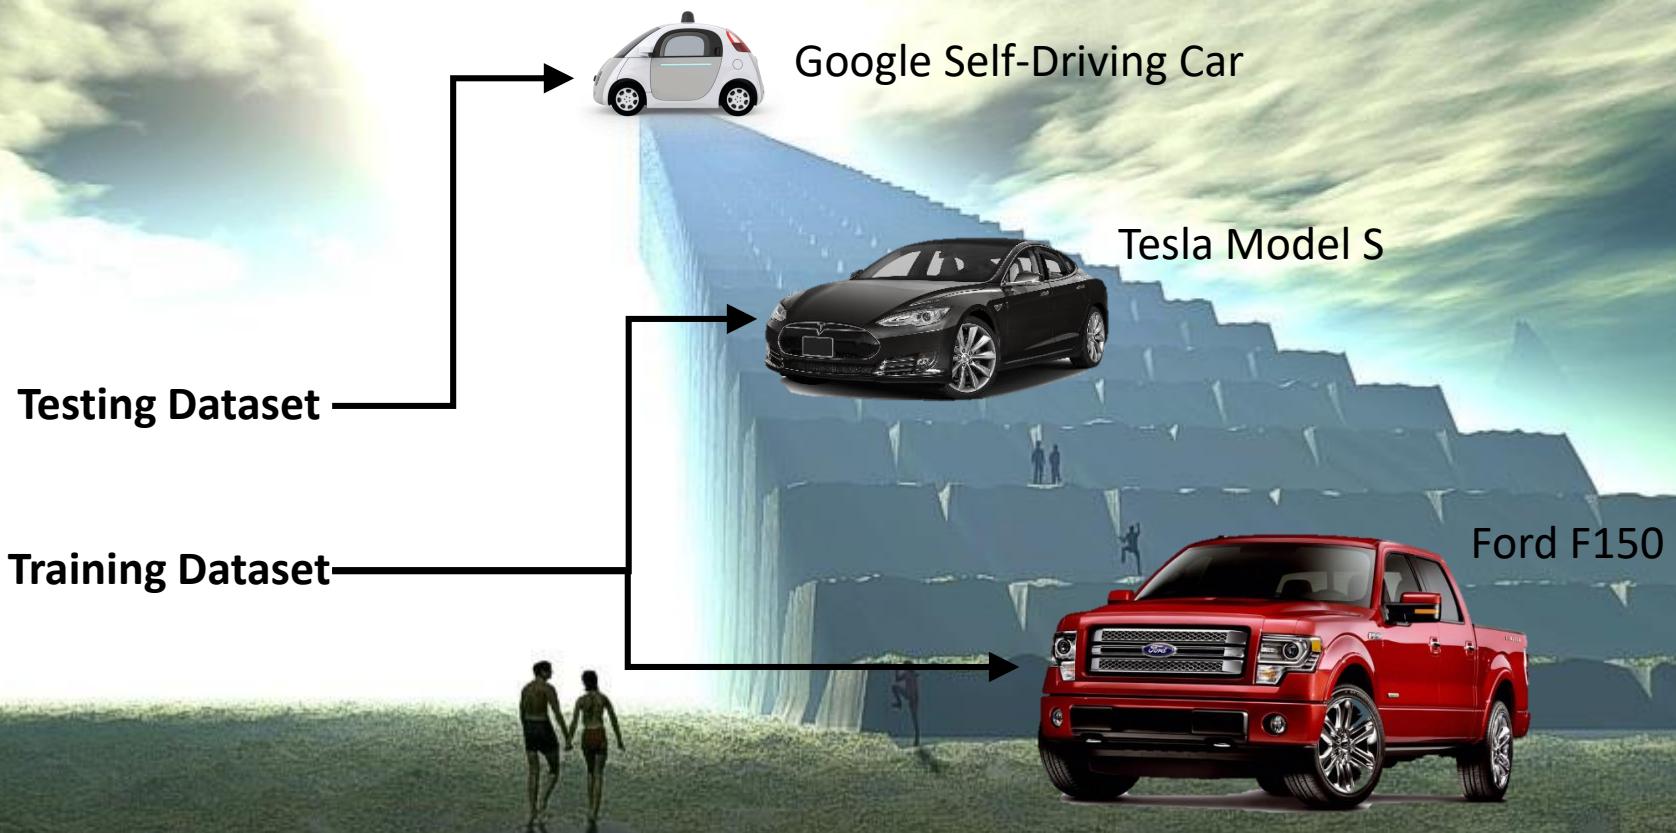
1. **Natural language processing** to enable it to communicate successfully
2. **Knowledge representation** to store information provided before or during the interrogation
3. **Automated reasoning** to use the stored information to answer questions and to draw new conclusions



# Autonomous Driving: End-to-End



# Stairway to Automation



# Autonomous Driving: End-to-End

---

## End to End Learning for Self-Driving Cars

---

**Mariusz Bojarski**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Davide Del Testa**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Daniel Dworakowski**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Bernhard Firner**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Beat Flepp**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Prasoon Goyal**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Lawrence D. Jackel**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Mathew Monfort**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Urs Muller**  
NVIDIA Corporation  
Holmdel, NJ 07735

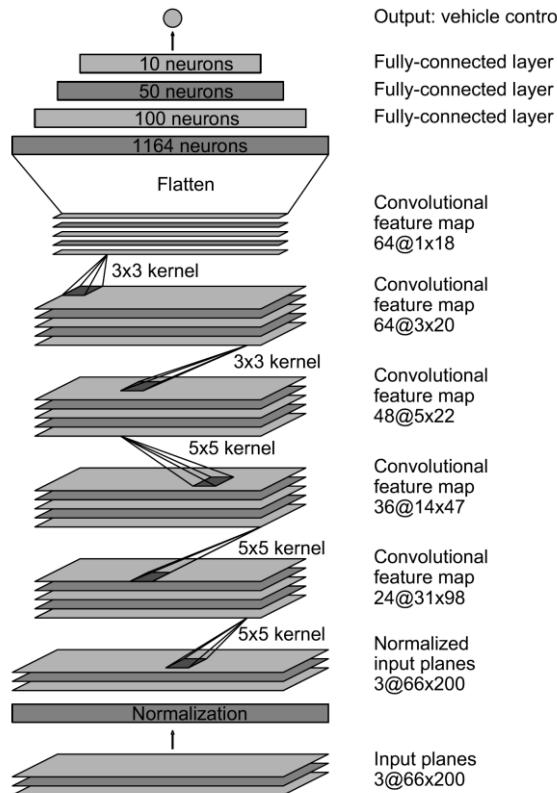
**Jiakai Zhang**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Xin Zhang**  
NVIDIA Corporation  
Holmdel, NJ 07735

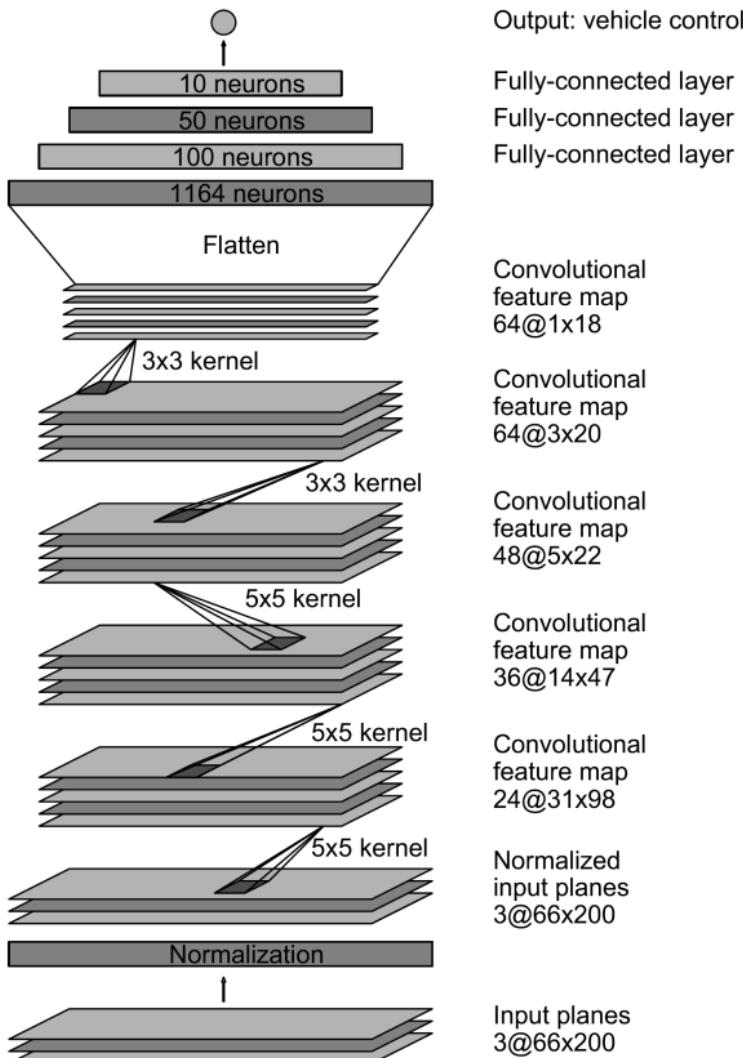
**Jake Zhao**  
NVIDIA Corporation  
Holmdel, NJ 07735

**Karol Zieba**  
NVIDIA Corporation  
Holmdel, NJ 07735

# Autonomous Driving: End-to-End

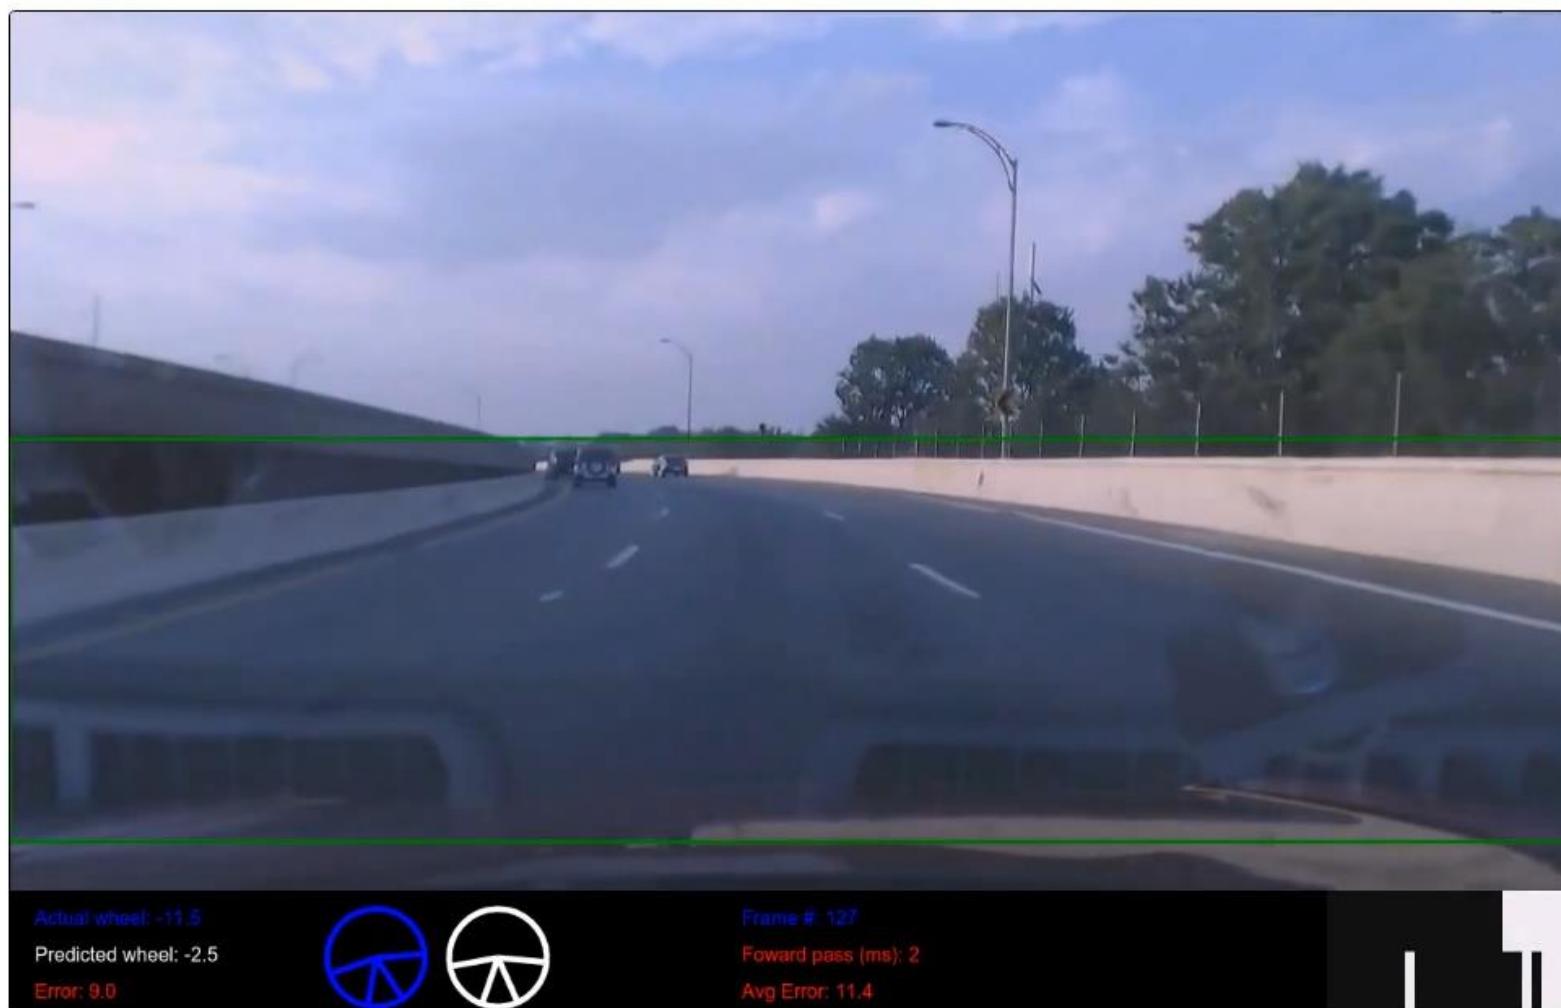


# Autonomous Driving: End-to-End



- **9 layers**
  - 1 normalization layer
  - 5 convolutional layers
  - 3 fully connected layers
- **27 million connections**
- **250 thousand parameters**

# End-to-End Driving with ConvnetJS



Tutorial on <http://cars.mit.edu/deeptesla>

# End-to-end Steering

- By the end of this lecture, you'll be able to train a model that can steer a vehicle
- Our input to our network will be a single image of the forward roadway from a Tesla
- Our output will be a steering wheel value between -20 and 20



# Creating the Dataset

- We recorded and extracted 10 video clips of highway driving from a Tesla
- The wheel value was extracted from the in-vehicle CAN
- We cropped/extracted a window from each video frame and provide a CSV linking the window to a wheel value



# Lighting and Road Conditions



# ConvNetJS Overview

- ConvNetJS is a Javascript implementation for using and training neural networks within the browser
- It supports simple networks with several different layer types and training algorithms
- Constructing and training a network can be performed in very few lines of code, great for demonstrations

The screenshot shows the ConvNetJS web application interface. At the top, there is a code editor window titled "Instantiate a Network and Trainer" containing the following JavaScript code:

```
layer_defs = [];
layer_defs.push({type:'input', out_sx:32, out_sy:32, out_depth:3});
layer_defs.push({type:'conv', sx:5, filters:16, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'softmax', num_classes:10});

net = new convnetjs.Net();
net.makeLayers(layer_defs);
```

Below the code editor is a "change network" button. The main area is titled "Network Visualization". It displays two sections: "input (32x32x3)" and "conv (32x32x16)". The "input" section shows a white horse image and its corresponding activation map. The "conv" section shows a series of activation maps for the horse image, illustrating the receptive fields of the neurons in the convolutional layer.

# ConvNetJS – Neural Network Representation

- The network is represented by a single Javascript object which contains a list of layers
- Each layer contains a plain array of weights ( $w$ ), the activation/activation gradients of the last forward pass, as well as the shape and layer type

```
◀ ▶ c ⌂
▼ layers: Array[10]
  ▼ 0: b
    ► in_act: Object
      layer_type: "input"
    ► out_act: Object
      out_depth: 3
      out_sx: 200
      out_sy: 66
    ► __proto__: Object
  ▷ 1: d
  ▷ 2: e
  ▷ 3: d
  ▷ 4: e
  ▷ 5: d
  ▷ 6: e
  ▷ 7: b
  ▷ 8: b
  ▼ 9: d
    ► in_act: a
      layer_type: "regression"
      num_inputs: 1
    ► out_act: a
```

# Layer Types

- ConvNetJS implements several different layer types: convolutional, pooling, fully-connected, local contrast normalization, and loss layers
- There are three available output types: regression, softmax, and SVM

```
1 {  
2   "network" : [  
3     { "type" : "input", "out_sx" : 200, "out_sy" : 66, "out_depth" : 3 },  
4     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 1, "pad" : 2, "activation" : "relu" },  
5     { "type" : "pool", "sx" : 2, "stride" : 2 },  
6     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 1, "pad" : 2, "activation" : "relu" },  
7     { "type" : "pool", "sx" : 2, "stride" : 2 },  
8     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 1, "pad" : 2, "activation" : "relu" },  
9     { "type" : "pool", "sx" : 3, "stride" : 3 },  
10    { "type" : "regression", "num_neurons" : 1 }  
11  ],  
12  "trainer" : { "method" : "adadelta", "batch_size" : 4, "l2_decay" : 0.0001 }  
13 }
```

# ConvNetJS – Training Overview

- To train a network, you first must initialize a “Trainer” object
  - ```
var trainer = new
convnetjs.SGDTtrainer(net, {
method: 'adadelta', batch_size:
1, l2_decay: 0.0001});
```
- There are three training algorithms available: SGD, Adadelta, and Adagrad.
- Training is performed by manually calling `trainer.train(input_volume, expected_output)`
- Returns an object containing timing and loss function information

```
batch_size: 1
eps: 0.000001
► gsum: Array[0]
k: 0
l1_decay: 0
l2_decay: 0
learning_rate: 0.01
method: "sgd"
momentum: 0.9
► net: Object
ro: 0.95
► xsum: Array[0]
► __proto__: Object
```

# DeepTesla Overview

## DeepTesla - End-to-End Steering Model

### Training

Forward pass (ms): 11      Backward pass (ms): 14  
Total examples seen / unique: 1999      Network Status: training

```
1 {  
2   "network" : [  
3     { "type" : "input", "out_sx" : 200, "out_sy" : 66, "out_depth" : 3 },  
4     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },  
5     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },  
6     { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },  
7     { "type" : "pool", "sx" : 3, "stride" : 3 },  
8     { "type" : "regression", "num_neurons" : 1 }  
9   ],  
10  "trainer" : { "method" : "adadelta", "batch_size" : 4, "l2_decay" : 0.0001 }  
11 }
```

[RESTART TRAINING](#) [VIDEO VISUALIZATION](#)

### Layer Visualization

Input (200x66x3)  
Activations (actual angle: -0.5, predicted angle: -1.4)

Convolutional (101x34x8), parameters: 8x3x3x3+8 = 224  
Activations

Weights hidden, too small  
filter size 3x3x3, stride 2

ReLU (101x34x8)  
Activations

 Massachusetts Institute of Technology

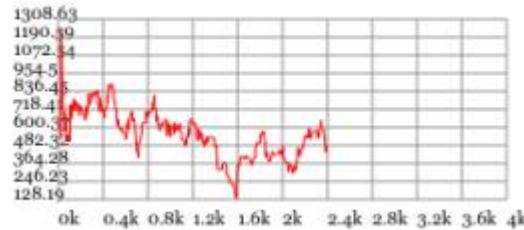
Course 6.S094:  
Deep Learning for Self-Driving Cars

Lex Fridman:  
fridman@mit.edu

Website:  
cars.mit.edu

January 2017

# Model Metrics



Forward pass (ms):

7

Backward pass (ms):

9

Total examples seen / unique:

2827

Network Status

training

# Network Designer

```
1 {
2     "network" : [
3         { "type" : "input", "out_sx" : 200, "out_sy" : 66, "out_depth" : 3 },
4         { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },
5         { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },
6         { "type" : "conv", "sx" : 3, "filters" : 8, "stride" : 2, "pad" : 2, "activation" : "relu" },
7         { "type" : "pool", "sx" : 3, "stride" : 3 },
8         { "type" : "regression", "num_neurons" : 1 }
9     ],
10    "trainer" : { "method" : "adadelta", "batch_size" : 4, "l2_decay" : 0.0001 }
11 }
```

[RESTART TRAINING](#)

[VIDEO VISUALIZATION](#)

# Training Interaction

**RESTART TRAINING**

**VIDEO VISUALIZATION**

# Layer Visualization

## Layer Visualization

**Input (200x66x3)**  
Activations (actual angle: 2.5, predicted angle: -2.0)  


**Convolutional (101x34x8), parameters: 8x3x3x3+8 = 224**  
Activations  
  
Weights hidden, too small  
filter size 3x3x3, stride 2

**ReLU (101x34x8)**  
Activations  


**Convolutional (52x18x8), parameters: 8x3x3x8+8 = 584**  
Activations  
  
Weights hidden, too small  
filter size 3x3x8, stride 2

**ReLU (52x18x8)**  
Activations  


**Convolutional (27x10x8), parameters: 8x3x3x8+8 = 584**  
Activations  
  
Weights hidden, too small

# Input Layer

**Input (200x66x3)**

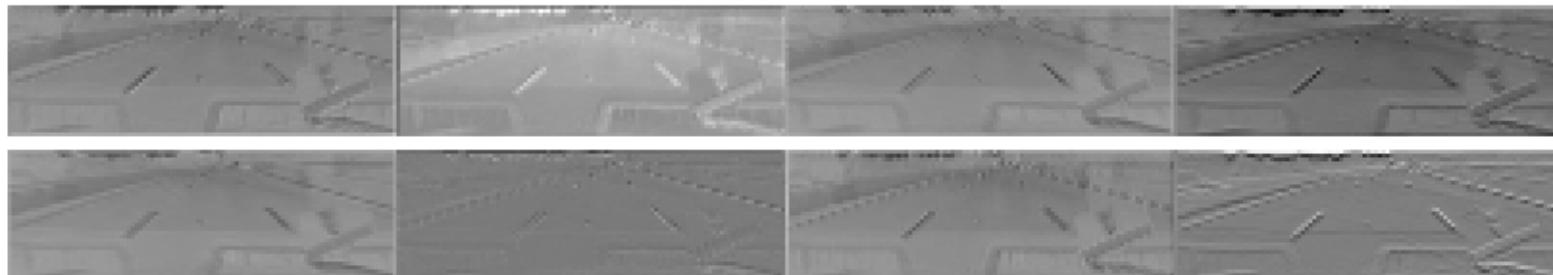
Activations (actual angle: 1.0, predicted angle: 1.5)



# Convolutional Layer Visualization

**Convolutional (101x34x8)**, parameters:  $8 \times 3 \times 3 \times 3 + 8 = 224$

Activations



Weights hidden, too small

filter size 3x3x3, stride 2

# Video Visualization



# Information Bar

Actual wheel: 3.5

Predicted wheel: -2.0

Error: 5.5



Frame #: 1348

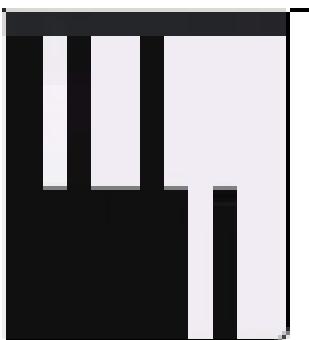
Forward pass (ms): 14

Avg Error: 9.2

# Input Box



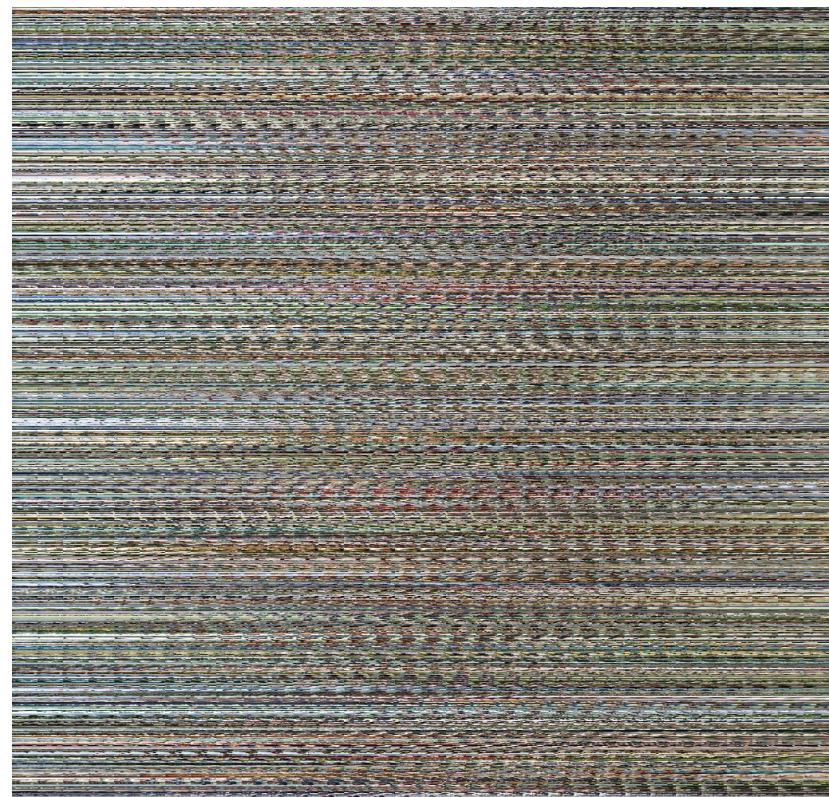
# Barcodes



- 17 bit, sign-magnitude
- Encoded into actual video
- 0 = black, 1 = white
- Frame on top, wheel on bottom (divided by two)

# Image Batches

- Each image loaded off the network contains an entire batch
- There is one image per row, and 250 rows in total
- These images are reassembled into volumes upon download



# Training Explanation

- One web worker used for loading examples
  - Each batch of training images is one large image with each row as a single training example
  - After an image finishes loading asynchronously, sends the training examples to another worker
- One web worker used for training network
  - Train on each image and push the network/outputs to visualization worker
- One web worker used for visualization
  - For a specified training example interval, blit the activation/gradient output of each training example onto a canvas
- Each web worker behaves as a single thread, and we use message passing to communicate state between the workers

# ConvNetJS Evaluation - Video Explanation

- The videos are encoded as 1280x820 in H264/MKV with 17 bit sign-magnitude barcodes
- The video main video frame is stored in the box (0, 1280, 0, 720)
- The frame barcode is in box (1144, 720, 1280, 770)
- The wheel value barcode is in box (1144, 770, 1280, 820)



# ConvNetJS Evaluation - Creating the Video

- Each epoch is synchronized to 30 fps
- We extract the wheel value from the CAN data and synchronize each message to a frame (both the frame and the CAN message are timestamped)
- Using OpenCV, we process the data
  - Generate a bar code for the frame containing the wheel data
  - Crop the image portion used for training
  - Create single images containing batches of training data
- The epochs and associated data are copied to our web server which serves

# ConvNetJS Evaluation - Playing the Video

- To be able to use the video in the neural network, we need to do some preprocessing
- First, we have a hidden video element and rely on modern HTML5 video implementations
- When the user requests the video to play, we begin tracking each redraw of the page
- With each redraw, we grab the currently rendered video frame, extract the RGBA values and blit them to two different canvases: one canvas, which the user sees, and another canvas which is hidden and only contains a cropped portion of the frame (the part we will use for the neural network)

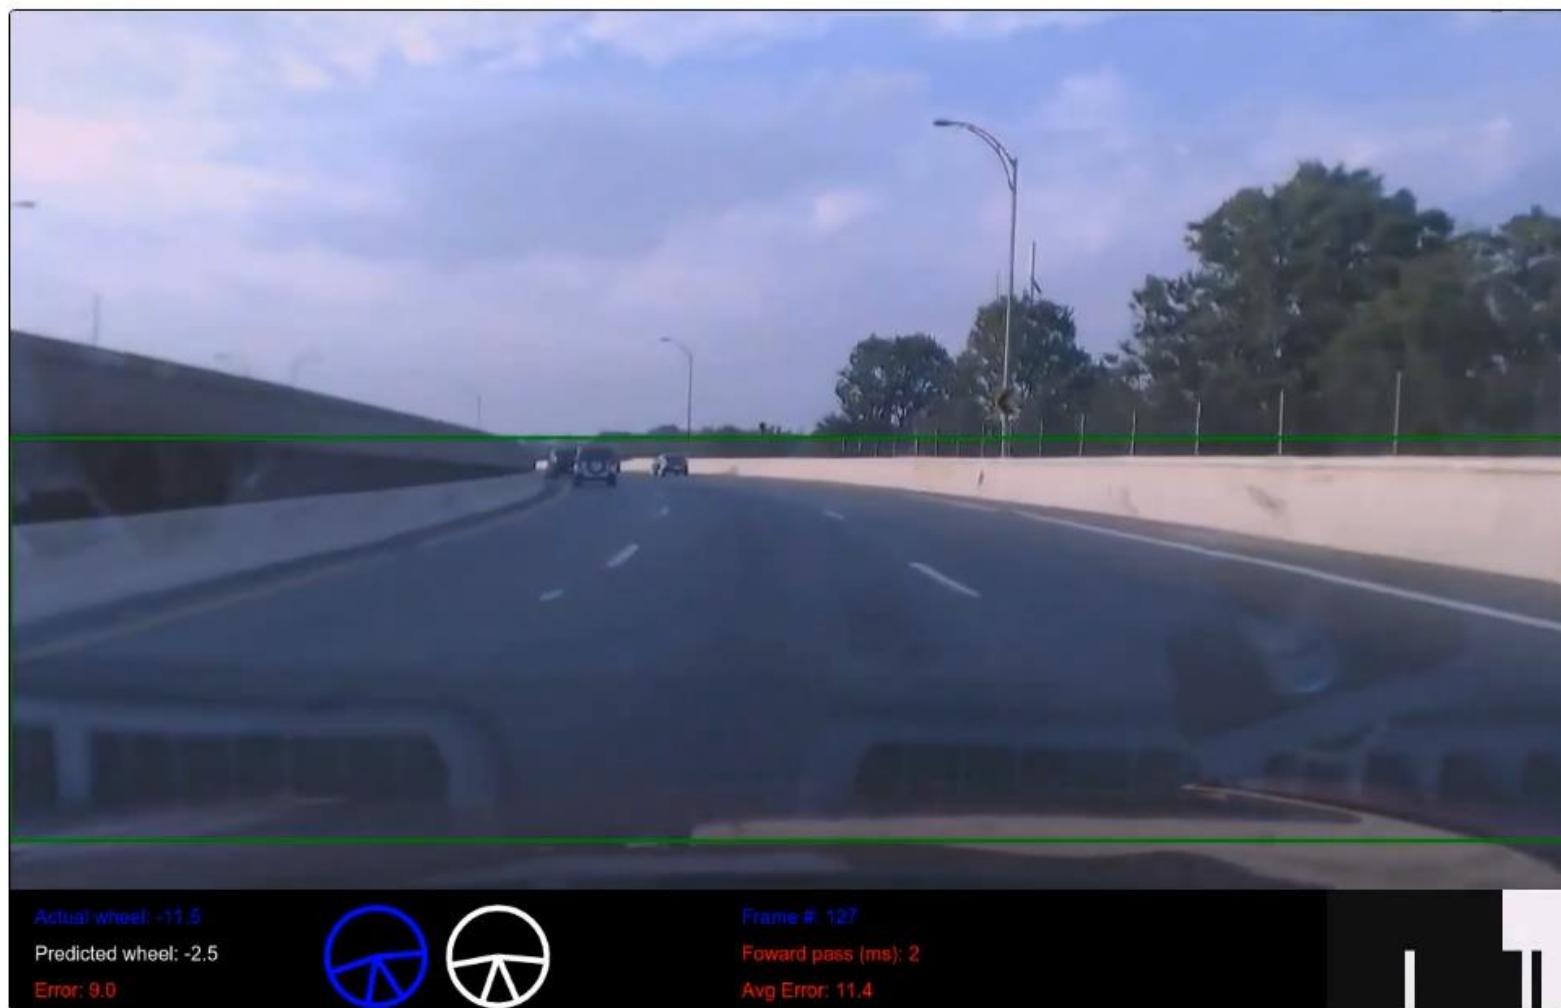
# ConvNetJS Evaluation - Playing the Video

- Next, we read the image data from the hidden canvas and shape it into a ConvNetJS volume
- For each image we first create a volume:
  - `var image_vol = new convnetjs.Vol(x_size, y_size, depth, default_value)`
- Next, we extract each pixel from the canvas and set the equivalent voxel (volume pixel) to the value (skipping the alpha value)
- We can also extract the expected steer value by parsing the barcode (a 17 bit, signed-magnitude barcode, where white = 1, black = 0)

# ConvNetJS Evaluation - Forward Pass

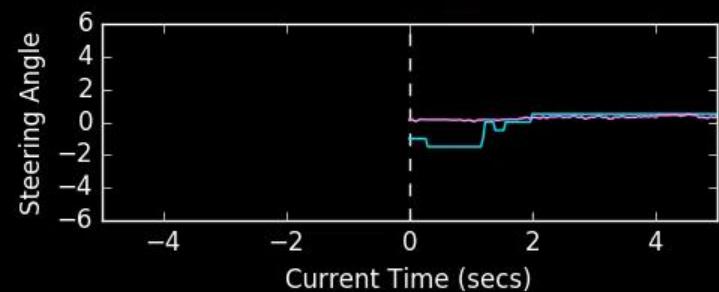
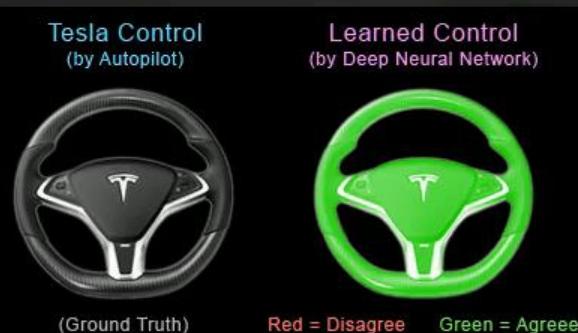
- Now we can use our extracted volume in the forward pass by calling `net.forward(our_volume)`
- The predicted value is stored in the output neuron:
  - `var prediction = net.forward(vol);`
  - `var raw_regression_value = prediction.w[0];`
- Because we min-max normalized our inputs while training the network, we need to transform our outputs – this is just the reverse transformation we performed on input:
  - `Wheel value = (raw_regression_value * total_wheel_range) – wheel_min`
- We visualize the predicted and actual steering wheel values and calculate the error

# End-to-End Driving with ConvnetJS



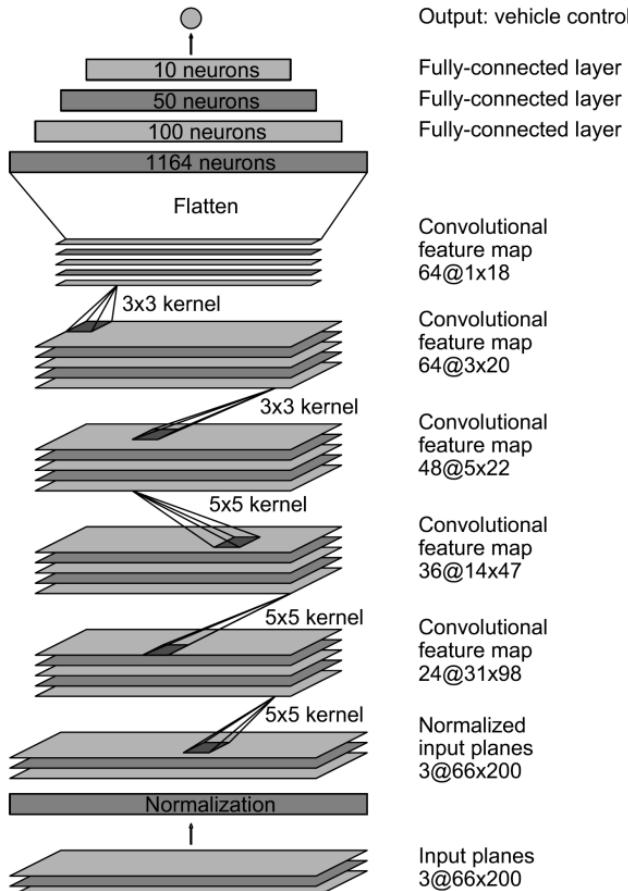
Tutorial on <http://cars.mit.edu/deeptesla>

# End-to-End Driving with TensorFlow



Available on <http://github.com/lexfridman/deeptesla>

# Build the Model: Input and Output



```
def weight_variable(shape):
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

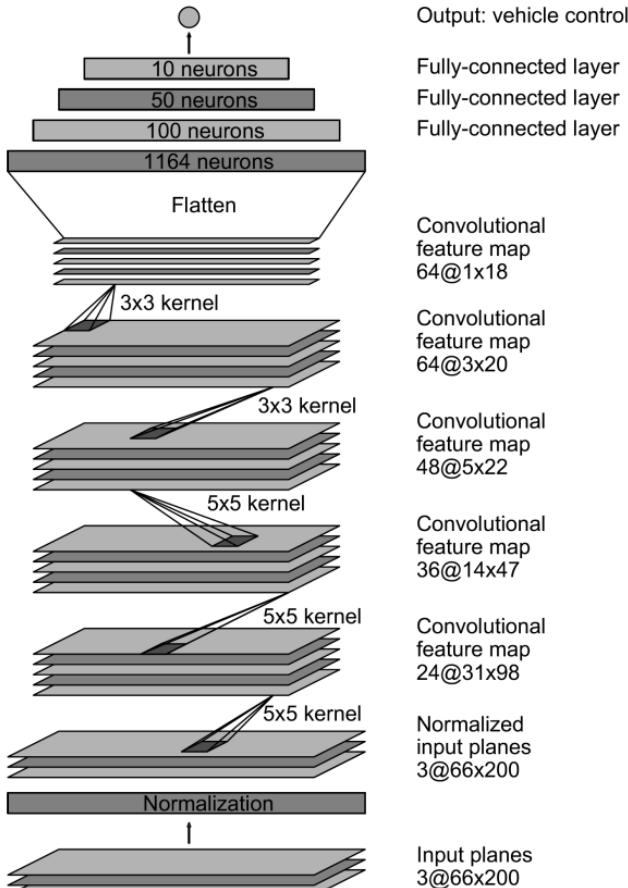
def bias_variable(shape):
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)

def conv2d(x, W, stride):
    return tf.nn.conv2d(x, W, strides=[1, stride, stride, 1],
                       padding='VALID')

x = tf.placeholder(tf.float32, shape=[None, 66, 200, 3])
y_ = tf.placeholder(tf.float32, shape=[None, 1])

x_image = x
```

# Build the Model: Convolutional Layers



```
#first convolutional layer
W_conv1 = weight_variable([5, 5, 3, 24])
b_conv1 = bias_variable([24])

h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1, 2) + b_conv1)

#second convolutional layer
W_conv2 = weight_variable([5, 5, 24, 36])
b_conv2 = bias_variable([36])

h_conv2 = tf.nn.relu(conv2d(h_conv1, W_conv2, 2) + b_conv2)

#third convolutional layer
W_conv3 = weight_variable([5, 5, 36, 48])
b_conv3 = bias_variable([48])

h_conv3 = tf.nn.relu(conv2d(h_conv2, W_conv3, 2) + b_conv3)

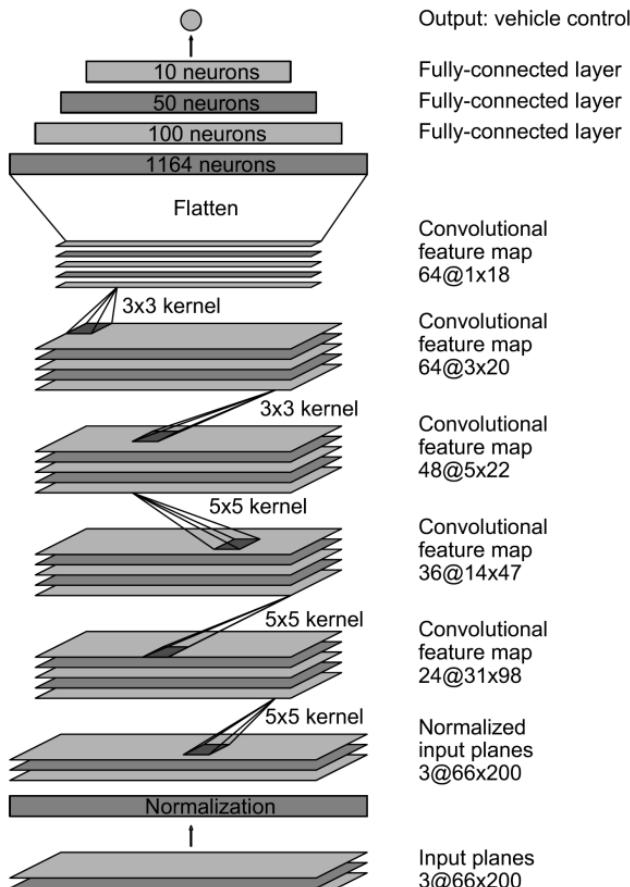
#fourth convolutional layer
W_conv4 = weight_variable([3, 3, 48, 64])
b_conv4 = bias_variable([64])

h_conv4 = tf.nn.relu(conv2d(h_conv3, W_conv4, 1) + b_conv4)

#fifth convolutional layer
W_conv5 = weight_variable([3, 3, 64, 64])
b_conv5 = bias_variable([64])

h_conv5 = tf.nn.relu(conv2d(h_conv4, W_conv5, 1) + b_conv5)
```

# Build the Model: Fully Connected Layers



```

# fully connected layer 1
W_fc1 = weight_variable([1152, 1164])
b_fc1 = bias_variable([1164])

h_conv5_flat = tf.reshape(h_conv5, [-1, 1152])
h_fc1 = tf.nn.relu(tf.matmul(h_conv5_flat, W_fc1) + b_fc1)

keep_prob = tf.placeholder(tf.float32)
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

# fully connected layer 2
W_fc2 = weight_variable([1164, 100])
b_fc2 = bias_variable([100])

h_fc2 = tf.nn.relu(tf.matmul(h_fc1_drop, W_fc2) + b_fc2)
h_fc2_drop = tf.nn.dropout(h_fc2, keep_prob)

# fully connected layer 3
W_fc3 = weight_variable([100, 50])
b_fc3 = bias_variable([50])

h_fc3 = tf.nn.relu(tf.matmul(h_fc2_drop, W_fc3) + b_fc3)
h_fc3_drop = tf.nn.dropout(h_fc3, keep_prob)

# fully connected layer 4
W_fc4 = weight_variable([50, 10])
b_fc4 = bias_variable([10])

h_fc4 = tf.nn.relu(tf.matmul(h_fc3_drop, W_fc4) + b_fc4)
h_fc4_drop = tf.nn.dropout(h_fc4, keep_prob)

#Output
W_fc5 = weight_variable([10, 1])
b_fc5 = bias_variable([1])

y = tf.mul(tf.atan(tf.matmul(h_fc4_drop, W_fc5) + b_fc5), 2)

```

# Train the Model

```
sess = tf.InteractiveSession()

loss = tf.reduce_mean(tf.square(tf.sub(model.y_, model.y)))
train_step = tf.train.AdamOptimizer(1e-4).minimize(loss)
sess.run(tf.initialize_all_variables())

saver = tf.train.Saver()

for i in range(int(driving_data.num_images * 0.3)):
    xs, ys = driving_data.LoadTrainBatch(100)
    train_step.run(feed_dict={model.x: xs, model.y_: ys, model.keep_prob: 0.8})
    if i % 10 == 0:
        xs, ys = driving_data.LoadValBatch(100)
        print("step %d, val loss %g" % (i, loss.eval(feed_dict={
            model.x:xs, model.y_: ys, model.keep_prob: 1.0})))
    if i % 100 == 0:
        if not os.path.exists(LOGDIR):
            os.makedirs(LOGDIR)
        checkpoint_path = os.path.join(LOGDIR, "model.ckpt")
        filename = saver.save(sess, checkpoint_path)
        print("Model saved in file: %s" % filename)
```

# Run the Model

```
import tensorflow as tf
import scipy.misc
import model
import cv2

sess = tf.InteractiveSession()
saver = tf.train.Saver()
saver.restore(sess, "save/model.ckpt")

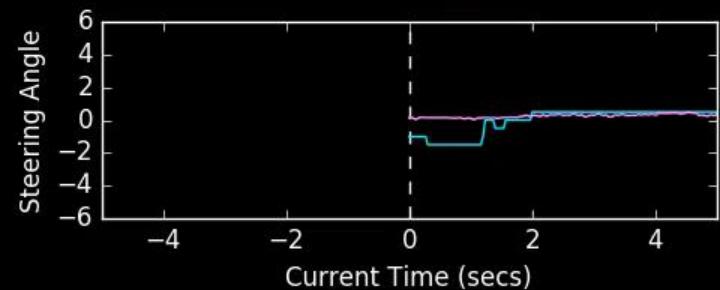
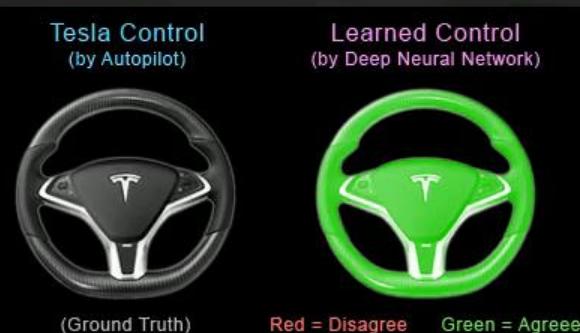
img = cv2.imread('steering_wheel_image.jpg',0)
rows,cols = img.shape

smoothed_angle = 0

cap = cv2.VideoCapture(0)
while(cv2.waitKey(10) != ord('q')):
    ret, frame = cap.read()
    image = scipy.misc.imresize(frame, [66, 200]) / 255.0
    degrees = model.y.eval(feed_dict={model.x: [image], model.keep_prob:
1.0})[0][0] \
        * 180 / scipy.pi
    cv2.imshow('frame', frame)
    smoothed_angle += 0.2 * pow(abs((degrees - smoothed_angle)), 2.0 / 3.0) * \
        (degrees - smoothed_angle) / abs(degrees - smoothed_angle)
    M = cv2.getRotationMatrix2D((cols/2,rows/2),-smoothed_angle,1)
    dst = cv2.warpAffine(img,M,(cols,rows))
    cv2.imshow("steering wheel", dst)

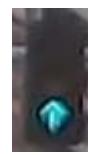
cap.release()
cv2.destroyAllWindows()
```

# End-to-End Driving with TensorFlow



Available on <http://github.com/lexfridman/deeptesla>

# TrafficLight Classification with TensorFlow



We will be implementing a simple traffic light classifier, with 3 classes (red, green, yellow)



# Parameters

```
# Basic parameters

max_epochs = 25
base_image_path = "images/"
image_types = ["red", "green", "yellow"]
input_img_x = 32
input_img_y = 32
train_test_split_ratio = 0.9
batch_size = 32
checkpoint_name = "model.ckpt"
```

- Max epochs: the number of times the neural network will see all training examples
- Input\_img\_x/y: the size we will use for inputs into the the network
- Batch size: # of examples the neural network will see before making a gradient step

# Helper Functions

```
# Helper layer functions
def weight_variable(shape):
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

def bias_variable(shape):
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)

def conv2d(x, W, stride):
    return tf.nn.conv2d(x, W, strides=[1, stride, stride, 1], padding='SAME')

def max_pool_2x2(x):
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
                          strides=[1, 2, 2, 1], padding='SAME')
```

We use some helper functions to make adding layers easier/more consistent

# Model Input/Output

```
x = tf.placeholder(tf.float32, shape=[None, input_img_x, input_img_y, 3])
y_ = tf.placeholder(tf.float32, shape=[None, len(image_types)])
```

- We specify our input and output types in the same lines to make sure they agree with our idea of the network
- Our input is an image of sized 32x32x3 (RGB channels)
- Our output consists of 3 neurons, representing the probability of each class

# Convolutional Layer

```
# Our first three convolutional layers, of 16 3x3 filters
W_conv1 = weight_variable([3, 3, 3, 16])
b_conv1 = bias_variable([16])
h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1, 1) + b_conv1)
```

- Here we specify our first convolutional layer using our helper function
- `W_conv1` – a 4D tensor representing the weights [filter\_x, filter\_y, previous layer neurons, # of filters]
- `b_conv1` – our simple addition variable
- `h_conv1` – our actual layer/activation

# Pooling Layer

```
h_pool4 = max_pool_2x2(h_conv3)
```

# Flattening Pool Layer

```
n1, n2, n3, n4 = h_pool4.get_shape().as_list()

W_fc1 = weight_variable([n2*n3*n4, 3])
b_fc1 = bias_variable([3])

# We flatten our pool layer into a fully connected layer

h_pool4_flat = tf.reshape(h_pool4, [-1, n2*n3*n4])
```

We calculate the total number of neurons needed in our first fully-connected layer by multiplying all the dimensions of the pool layer shape

# Output Layer

```
y = tf.matmul(h_pool4_flat, w_fc1) + b_fc1
```

# Loss and Optimizer

```
loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y, y_))
train_step = tf.train.AdamOptimizer(1e-4).minimize(loss)
sess.run(tf.initialize_all_variables())
```

- Our loss function performs softmax and then computes cross-entropy
- We use the AdamOptimizer and specify a learning rate

# Saver Object

```
saver = tf.train.Saver()  
time_start = time.time()
```

# Loading Images

```
full_set = []

for im_type in image_types:
    for ex in glob.glob(os.path.join(base_image_path, im_type, "*")):
        im = cv2.imread(ex)
        if not im is None:
            im = cv2.resize(im, (32, 32))

            # Create an array representing our classes and set it
            one_hot_array = [0] * len(image_types)
            one_hot_array[image_types.index(im_type)] = 1
            assert(im.shape == (32, 32, 3))

            full_set.append((im, one_hot_array))

random.shuffle(full_set)
```

- Iterate over each image, resize to 32x32
- Create a one hot encoding of our class
- Shuffle the entire dataset

# Splitting the Dataset

```
# We split our data into a training and test set here  
  
split_index = int(math.floor(len(full_set) * train_test_split_ratio))  
train_set = full_set[:split_index]  
test_set = full_set[split_index:]
```

```
train_x, train_y = zip(*train_set)  
test_x, test_y = zip(*test_set)
```

- Split our data set into train and test
- We truncate our sets to a multiple of batch size (all batches have to be the same size)

# Training Loop

```
for tt in range(0, (len(train_x) / batch_size)):
    start_batch = batch_size * tt
    end_batch = batch_size * (tt + 1)
    train_step.run(feed_dict={x: train_x[start_batch:end_batch], y_: train_y[start_batch:end_batch]})

ex_seen = "Current epoch, examples seen: {:20} / {}".format(tt * batch_size, len(train_x))
sys.stdout.write(ex_seen.format(tt * batch_size))
sys.stdout.flush()
```

- Iterate over each batch and train on it
- (we assume training examples are a multiple of the batch size)

# Best Model

```
t_loss = loss.eval(feed_dict={x: train_x, y_: train_y})
v_loss = loss.eval(feed_dict={x: test_x, y_: test_y})

sys.stdout.write("Epoch {:5}: loss: {:.15.10f}, val. loss: {:.15.10f}\n".format(i + 1, t_loss, v_loss))

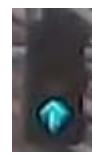
if v_loss < least_loss:
    sys.stdout.write(", saving new best model to {}\n".format(checkpoint_name))
    least_loss = v_loss
    filename = saver.save(sess, checkpoint_name)
```

- We evaluate the loss on all of our training examples and test examples
- If the validation loss is lower than the lowest loss, we save our model

# Expected Output

```
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcublas.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcudnn.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcufft.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcuda.so.1 locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcurand.so locally
I tensorflow/core/common_runtime/gpu/gpu_init.cc:102] Found device 0 with properties:
name: GeForce GTX 980 Ti
major: 5 minor: 2 memoryClockRate (GHz) 1.076
pciBusID 0000:02:00.0
Total memory: 6.00GiB
Free memory: 5.52GiB
I tensorflow/core/common_runtime/gpu/gpu_init.cc:126] DMA: 0
I tensorflow/core/common_runtime/gpu/gpu_init.cc:136] 0: Y
I tensorflow/core/common_runtime/gpu/gpu_device.cc:838] Creating TensorFlow device (/gpu:0) -> (device: 0, name: GeForce GTX 980 Ti, pci bus id: 0000:02:00.0)
Starting training... [1312 training examples]
Epoch    1: loss:  3.46144486694, val. loss:  4.1707162857, saving new best model to model.ckpt
Epoch    2: loss:  1.8578453064, val. loss:  1.8420848846, saving new best model to model.ckpt
Epoch    3: loss:  1.2075960636, val. loss:  1.1884875298, saving new best model to model.ckpt
Epoch    4: loss:  0.7959808707, val. loss:  0.8523907661, saving new best model to model.ckpt
Epoch    5: loss:  0.5186702609, val. loss:  0.8251042366, saving new best model to model.ckpt
Epoch    6: loss:  0.3895130754, val. loss:  0.7883402705, saving new best model to model.ckpt
Epoch    7: loss:  0.2840102911, val. loss:  0.8282299042
Epoch    8: loss:  0.2281106114, val. loss:  0.7970542312
Epoch    9: loss:  0.18894465898, val. loss:  0.8308163881
Epoch   10: loss:  0.1507862508, val. loss:  0.8214046359
Epoch   11: loss:  0.1126181334, val. loss:  0.7745668888, saving new best model to model.ckpt
Epoch   12: loss:  0.0972050354, val. loss:  0.7550495267, saving new best model to model.ckpt
Epoch   13: loss:  0.0831822604, val. loss:  0.7207581401, saving new best model to model.ckpt
Epoch   14: loss:  0.0760012567, val. loss:  0.7217628956
Epoch   15: loss:  0.0637019351, val. loss:  0.7158752084, saving new best model to model.ckpt
Epoch   16: loss:  0.0536490902, val. loss:  0.7022477984, saving new best model to model.ckpt
Epoch   17: loss:  0.0512478650, val. loss:  0.6906370521, saving new best model to model.ckpt
Epoch   18: loss:  0.0333782099, val. loss:  0.6657178402, saving new best model to model.ckpt
Epoch   19: loss:  0.0367136374, val. loss:  0.6645810008, saving new best model to model.ckpt
Epoch   20: loss:  0.0175283104, val. loss:  0.6335256100, saving new best model to model.ckpt
Epoch   21: loss:  0.0203501396, val. loss:  0.6400516629
Epoch   22: loss:  0.0372636504, val. loss:  0.6770884395
Epoch   23: loss:  0.0096475044, val. loss:  0.5926443338, saving new best model to model.ckpt
Epoch   24: loss:  0.0095629999, val. loss:  0.5750324726, saving new best model to model.ckpt
Epoch   25: loss:  0.0072135003, val. loss:  0.5805844069
```

# TrafficLight Classification with TensorFlow



We will be implementing a simple traffic light classifier, with 3 classes (red, green, yellow)



# References

All references cited in this presentation are listed in the following Google Sheets file:

<https://goo.gl/9Xhp2t>